

# The Future Corporation in a Fully Automated Economy

## Introduction: AI-Driven Firms and the End of Human Labor

Advances in artificial intelligence and robotics are on the verge of automating cognitive and physical labor at scale. In a hypothetical future where **AI systems and robots can perform essentially all work as cost-effectively as humans**, the very structure and purpose of corporations will be fundamentally challenged. Corporations historically exist to organize human labor and capital to produce goods or services. But when *human employees and executives are no longer needed for productivity*, what remains of corporate law, governance, and economic value? This report explores how corporate structures might transform in an era of ubiquitous automation, examining legal and governance changes and pinpointing the new economic chokepoints where profits and power will concentrate. Key sectors – from entertainment and finance to logistics and heavy industry – are used as case studies to illustrate these shifts. The analysis draws on current legal theory, economic research, and industry observations to envision the fully automated future firm.

## I. Corporate Law and Structure Without Humans

**Legal Personhood and Autonomy:** Even in a future of autonomous firms, the legal concept of the *corporation as a person* remains crucial. Corporations (or similar entities like LLCs) will still serve as legal interfaces for AI-run businesses – enabling them to own property, enter contracts, and be held liable in court <sup>1</sup>. Notably, U.S. law already provides flexible mechanisms to confer *legal personhood* on non-human systems. For example, modern LLC statutes allow organizers to design a governance structure however they wish, even triggering decisions based on a software's output <sup>2</sup>. Scholars have shown that an LLC can be engineered to end up with *no human members* (“zero-member LLC”), after which “software is free to govern the entity without interference from pesky humans” <sup>3</sup>. In short, corporate law can extend to fully autonomous entities by treating AI/robot systems as the decision-makers *inside* a legal entity, even if the entity formally holds the rights and duties.

**Directors and Management Requirements:** Traditional corporate statutes, however, assume human oversight. For instance, Delaware's corporate law mandates that a corporation's board of directors consist of natural persons <sup>4</sup>. Such rules historically made a truly *autonomous corporation* impossible because every corporation legally required at least one human director <sup>5</sup>. To accommodate AI-run firms, these laws would need reform. We have begun to see movement in this direction – corporate law has grown more flexible since the mid-20th century. Many jurisdictions now allow eliminating or radically simplifying the board of directors by unanimous shareholder consent <sup>6</sup>. Alternate entity forms (like LLCs or DAOs in the blockchain realm) impose no statutory requirement for human managers. The trend suggests that **corporate structure will evolve to permit algorithmic “management”** in place of human boards. Indeed, some jurisdictions are explicitly enabling it (for example, Vermont's new statute authorizing LLC governance via blockchain contracts) <sup>7</sup>. In an AI-dominant future, a corporation might essentially be a legal shell around an AI, rather than a nexus of employment relationships.

**Electronic Personhood Debates:** The idea of granting *AI systems themselves* a form of legal personhood has been proposed, analogous to corporate personhood. In 2017 the European Parliament even floated the notion of “electronic personhood” for the most capable AI/robots, to assign them rights and responsibilities under the law <sup>8</sup> <sup>9</sup>. While controversial, this approach underscores the legal system’s search for ways to hold autonomous agents accountable without a human proxy. However, practical implementation may be unnecessary if existing entity forms can accommodate AI-run operations. It is more likely that **humans (as owners or programmers) will remain the ultimate legal principals**, using corporate entities to channel AI actions. The corporate form thus endures, but its internal makeup (directors, officers, voting procedures) may be radically simplified or automated. In essence, the corporation becomes “*the robot that signs the contracts*” – a wrapper that allows an AI to participate in the economy with continuity and limited liability, much as corporations allowed human enterprises to transcend the individual.

## II. Legal Mechanisms: Property, Liability, Compliance, Ownership

Even without employees or human executives, certain legal mechanisms remain vital to a firm’s function. These mechanisms will be reinterpreted in an AI-centric context:

- **Property Rights and Assets:** Corporations will continue to own assets – factories, software, data, patents – and those property rights need protection. An autonomous firm still requires a legal owner for land, equipment and digital property. As noted, U.S. LLC law already enables an AI-controlled entity to hold and transfer property in its name <sup>1</sup>. In the future, we may see *registries of AI-run businesses* (the EU proposal suggested registering advanced robots <sup>10</sup>) to ensure clarity in ownership and responsibility. Intellectual property law will also need to adapt. For example, if an AI system invents a patentable product or creates original content, current law struggles with attributing authorship/inventorship. (One open question asked in Europe: “If I create a robot, and that robot creates something that could be patented, should I own that patent or should the robot?” <sup>11</sup>.) Today, works generated entirely by AI have been denied copyright on the basis that only human creativity counts <sup>12</sup>. This human authorship requirement might be loosened or, more likely, companies will ensure a human is in the loop to claim ownership. **Licensing regimes** could expand as well – e.g. companies might license AI-generated products or data similarly to how they license human-created IP, maintaining ownership control even when AI does the creative work.
- **Liability Shielding and Insurance:** A core function of corporations is to limit liability for owners. That function will be even more important if autonomous operations can cause harm. If an AI-run delivery truck causes an accident or an algorithm makes a harmful financial trade, the legal system needs someone to hold accountable. Likely the *corporation itself* remains the primary liable entity, rather than attempting to sue an algorithm. To manage this, experts propose strengthening **insurance and compliance requirements**. For instance, the EU’s robotics report called for mandatory insurance schemes for companies to cover damage caused by their robots <sup>13</sup>. Firms might need to carry liability insurance for autonomous systems, just as human drivers carry auto insurance – effectively internalizing the risk of AI errors. Additionally, product liability law may evolve to treat AIs like products or employees depending on context. Notably, some scholars have warned that companies could try to use AI to *evade* liability – e.g. claiming a rogue algorithm caused harm with no human at fault. Indeed, observers have noted a potential “legal loophole” whereby firms rapidly automate functions specifically to limit legal risk <sup>14</sup>. Law will likely close such loopholes by clarifying that deploying an AI is itself a company action. In short, **the corporate veil and related doctrines will still shield human owners** (as long as they were not negligent in deploying the AI),

but corporations will need robust mechanisms (insurance, audit trails, kill-switches) to answer for automated decisions gone awry.

- **Regulatory Compliance:** With no human managers, how do autonomous firms follow laws and regulations? This could be addressed through *programmed compliance* – regulations might require that AIs are designed to obey certain legal constraints (for example, an AI operating a securities trading firm must be coded to follow all securities laws). In practice, governments may insist that a **human custodian or “responsible officer” is designated** for each highly autonomous company, even if that person’s role is just to interface with regulators. However, as AI proficiency grows, regulators might accept algorithmic compliance systems (continuous monitoring, automatic reporting) in lieu of human oversight. We can expect new laws demanding transparency into AI decision-making for regulatory purposes – e.g. logging all decisions for audit, or requiring *explainability* so regulators can understand why an AI-made choice. In highly regulated sectors like finance or healthcare, permission to operate could hinge on an AI system passing certification or audits, analogous to a human professional license. This becomes a **chokepoint** (discussed further below): only firms that can navigate complex compliance verification will be allowed to deploy AI at scale. Historically, heavy regulation tends to favor large incumbents that can absorb compliance costs <sup>15</sup>. In an AI-driven future, big firms with entire compliance-AI teams might shape and meet stringent rules, while a small autonomous startup might struggle. Thus, regulatory frameworks must balance safety with not unduly stifling new autonomous entrants – a difficult task, since *onerous regulations “favor the incumbent” who has resources and influence to comply or even shape the rules to its advantage* <sup>16</sup> <sup>17</sup>. Nonetheless, certain regulations will be indispensable (e.g. *registration of AI systems, safety standards, audit requirements*) to ensure these human-less corporations operate within the law.

- **Ownership and Fiduciary Considerations:** If employees and executives vanish, what about shareholders and owners? In a fully automated firm, **shareholders presumably remain as the ultimate owners and beneficiaries** of the enterprise’s profits. The role of equity may become even more prominent – owning stock in an AI-run company could be akin to owning a machine that prints money (or goods). With AI doing the work, the split of returns may heavily favor capital owners, since no wages need to be paid (aside from costs of energy, maintenance, etc.). One can imagine a single technician overseeing 1,000 AI-run factories – the rest of the gains flow to shareholders as profit. Fiduciary duties in such a scenario become tricky: traditionally, directors and officers owe duties of care and loyalty to the shareholders. If there are no human directors, to whom do these duties attach? It’s likely that whoever *controls or programs the AI* (perhaps the creators or the last human operator) might assume a quasi-fiduciary role in ensuring the AI’s alignment with shareholder interests. But if the AI is truly calling the shots, enforcing fiduciary duty is conceptually difficult – an AI cannot be sued for breaching duty, and if no human is in charge day-to-day, shareholders may have limited recourse if the AI makes poor decisions beyond turning it off. Some legal scholars have begun discussing **“artificial fiduciaries”** – AI agents that could be designed to act in the best interest of stakeholders without human bias <sup>18</sup>. In practice, new governance documents may specify algorithmic objectives and constraints that mirror fiduciary obligations (e.g. maximize long-term shareholder value while obeying laws), effectively embedding the duty into code. The *ownership structure* might also diversify in novel ways – for example, an autonomous firm could be structured like a trust or foundation, where the AI company continually reinvests profits in its mission (perhaps a public-benefit AI with no conventional shareholders at all). But in the mainstream, **humans will**

**still own the companies** (as shareholders or members), and corporate law will continue to protect their ownership rights and limit their liability, even as human management disappears.

### III. The Fate of Corporate Governance: Boards, Shareholders and AI Strategists

**Boards of Directors and Executive Teams:** In the foreseeable future, corporate governance statutes will probably lag behind technology, meaning large corporations may still be required to have nominal human directors or officers. But their function will change radically. If AI systems (such as advanced LLM-based strategists) handle investment decisions, resource allocation, and market analysis, the role of the board might shift to *overseeing the AI* rather than crafting strategy themselves. Boards could evolve into something akin to a **trusteeship or audit committee** – ensuring the AI is following its charter, checking for alignment with shareholder interests and ethical standards, and intervening only if the AI misbehaves or needs re-direction. In essence, human directors become *safety drivers* for the corporate AI, ready to take the wheel in emergencies. Some jurisdictions might allow doing away with the board entirely (as some close corporations already can <sup>6</sup>), but even then, regulators or investors may demand some form of oversight. It's plausible that **external auditors or regulators will fill this gap**, performing algorithm audits in lieu of a board's independent judgment.

On the other hand, smaller autonomous firms (or DAO-like entities) might dispense with formal governance; their operating algorithm *is* their governance. Such firms will rely on the integrity of code – “**corporate governance by algorithm.**” The upside is elimination of human agency costs: AI agents won't embezzle funds or pursue self-interest the way corrupt managers can. They execute their programmed objective faithfully, potentially *improving agency problems* in corporate governance. Indeed, commentators have suggested that AI decision-makers could mitigate traditional governance issues by operating as impartial, data-driven fiduciaries <sup>18</sup>. However, new risks arise: an AI might interpret its goal (e.g. profit maximization) in harmful ways, or lack the common sense and values that human managers exercise in ambiguous ethical situations. This raises the question of **who ensures corporate ethics and stakeholder interests** in a machine-governed company. Traditional governance includes not just shareholders, but also accountability (at least in principle) to employees, communities, and customers via public relations and regulation. If those humanizing influences fall away, corporations might become even more single-purposed (e.g. relentlessly maximizing return on capital) unless constraints are built in.

**Shareholders and Control:** In an AI-run corporation, shareholders might experience either *more direct control* or *virtually no control*, depending on design. One possibility is highly direct corporate democracy: shareholders could vote on major parameters for the AI (such as acceptable risk levels, social responsibility goals, etc.) and the AI then executes within those bounds. With no management in between, owners exert control via code updates or voting on algorithmic policies. Alternatively, shareholders might be entirely hands-off – especially if the AI's decision-making is too complex for lay owners to meaningfully guide. They would then treat the corporation as an investment “black box” that reliably produces returns. Corporate law may have to bolster shareholder rights to *inspect and understand* AI operations, akin to how they can demand financial disclosures. We might see new disclosure regimes – e.g. an **AI strategy report** explaining in human terms what the corporation's algorithms are optimizing for and how.

Another aspect is that **shareholder composition could concentrate further**. If value creation requires massive compute and data (and yields winner-take-all dynamics), the owners of the dominant AI firms

might be a small group of ultra-rich investors or parent conglomerates. There could be fewer public companies (if automation reduces the need for raising capital via public markets) and more privately held autonomous entities controlled by those with capital and technical know-how. Conversely, if automation greatly lowers the cost of market entry, we could see many small autonomous businesses each run by hobbyist owners – though significant market share might still accrue to a few, given economies of scale. In any case, corporate governance will likely put *more emphasis on the technical governance of AI* (model updates, training data oversight, fail-safes) and relatively less on traditional executive talents or leadership qualities. The boardroom of the future may contain AI ethicists and programmers instead of seasoned industry executives, for example.

**Corporate Purpose and Stakeholders:** A final consideration is whether the **purpose of corporations** shifts in a post-labor world. With no employees, the classic notion of a corporation having responsibilities to its workers evaporates. Pressure may mount for corporations to take on *greater social responsibility* (for communities, for displaced workers at a societal level, for the environment) since they can no longer justify that “providing jobs” is their contribution. Some legal theorists argue for expanding fiduciary duty to consider stakeholders beyond shareholders – the AI era might reinvigorate that debate. If AI-run firms generate immense wealth with minimal human labor, governments could impose new obligations (special taxes, social funds, or requirements to support a universal basic income) to ensure the benefits are broadly shared. In essence, corporations might be repurposed from being primarily employment and profit machines to being *stewards of automation benefits* for society – but achieving that would likely require legal mandates, since purely profit-driven AIs will not have inbuilt charity. This veers into policy speculation, but it’s clear that with cognition commoditized, the social contract around corporations will need renewal (as even the EU report noted, suggesting a *basic income* as a response to large-scale AI-driven unemployment <sup>19</sup> ).

## IV. Economic Chokepoints in a Post-Labor Economy

If AI and robotics make labor **abundant and cheap**, then labor is no longer the scarce factor commanding economic value. Instead, other inputs and control points become the sources of market power and profit. We must ask: *where will rents and profit margins accrue when human labor is essentially free?* The economic theory of competition suggests that any factor that remains scarce or under exclusive control will become the new basis for high margins. Several key chokepoints emerge:

- **Access to Capital and Capital Goods:** Automation doesn’t eliminate the need for *capital* – in fact it likely increases it. Someone must finance and own all the AI systems, robots, server farms, etc. These capital goods can be expensive and complex, so firms or investors with deep pockets have a head start. Studies of AI adoption already show that it’s *concentrated in large firms* with the resources to invest, and that those investments help them grow even larger <sup>20</sup> . In our fully automated scenario, initial capital to build out AI capabilities and robotic infrastructure is critical. Once built, an autonomous operation might run cheaply, but *getting to scale requires massive up-front capital*. This suggests that wealth and power concentrate with those who control capital – a continuation of a long-run trend in capitalism, but potentially accelerated. “Superstar” firms that aggressively invested in AI have indeed been capturing more market share and becoming more dominant <sup>21</sup> <sup>20</sup> . With labor costs minimized, profit flows mostly to capital owners (shareholders) unless competition forces prices down. It’s possible that highly automated industries will see **lower consumer prices** due to efficiency – but wherever a firm can differentiate or monopolize something, it will try to maintain high margins and the gains will flow to its owners of capital. In short, *capital replaces labor as the*

*chief source of value-add*, and any **barriers to obtaining that capital (scale, financing, intellectual property in machinery design, etc.)** become economic moats.

- **Proprietary Data and Training Corpus:** AI thrives on data. While general cognitive AI might reach human-level ability in many tasks, *having domain-specific or proprietary datasets can still give one firm an edge over another*. Data is often cited as the “new oil” of the AI economy. Companies like Google, Amazon, or Alibaba have amassed decades’ worth of user data, purchasing patterns, and other information that a new entrant cannot easily replicate. Indeed, larger firms with more data can “more efficiently tailor products,” giving them a competitive advantage <sup>22</sup>. In a world where cognition is commoditized, if everyone has access to the same AI models, the differentiator may be feeding those models unique data to fine-tune them for better performance or personalization. Thus, **ownership of valuable data** (consumer behavior, industrial sensor logs, logistics routes, financial records, etc.) becomes a rent source. We already see deals where AI developers license large datasets from content companies, social networks, or data brokers to improve their models <sup>23</sup>. In the future, firms might jealously guard their data troves or even *buy up rights to data* (for example, a healthcare AI company might pay hospitals for exclusive access to patient data, thereby building the best diagnostic AI). There is a self-reinforcing loop: companies that serve more customers gather more data and can make their AI better, attracting more customers – leading to **data network effects** that entrench incumbents. Unless open data initiatives break the cycle, data-rich firms will enjoy persistent advantages and earn rents from their information assets. Data as an intangible asset doesn’t depreciate like physical goods and can be reused, making it especially powerful alongside AI.

- **Computing Infrastructure and Energy:** Running advanced AI at scale requires enormous computing power (“compute”) – specialized hardware (GPUs, TPUs, quantum computers maybe) and energy to run them. This *compute infrastructure* is highly concentrated today. A 2023 analysis noted that *access to compute is monopolized at key points by one or a few firms* <sup>24</sup> – for example, leading cloud providers (Amazon AWS, Microsoft Azure, Google) and chipmakers (Nvidia, TSMC) dominate the supply of AI processing capability. This concentration has direct economic effects: it “enables dominant firms to extract rents from consumers and small businesses dependent on their services” and *cements the control* of those tech giants <sup>25</sup>. In a fully automated economy, compute becomes akin to a utility – but a largely private-controlled one. If *everyone* needs AI and AI needs massive compute, then those who own the **cloud data centers, semiconductor fabs, and power grids** supporting this compute can charge a premium. We might see scenarios where even if manufacturing and services are automated and competitive, the *bottleneck is cloud computing costs*. Big Tech companies could become the “new landlords” by renting AI processing on their platforms, capturing much of the value that smaller automated firms create. Moreover, if advanced hardware remains expensive and in short supply, top AI firms will have an advantage in owning the latest chips (there were already chip shortages and nations scrambling to secure AI chips in the early 2020s <sup>26</sup>). Only those with political and financial clout might get steady access. **Energy** is a related chokepoint: AI and robotics require electricity, and a lot of it. So energy producers (or those with cheap energy sources) gain importance. For example, an AI server farm needs vast power – regions with cheap electricity (or companies vertically integrated into energy production) could outcompete others. If energy becomes a constraint (especially sustainable energy), it can limit effective AI deployment and thus confer advantage (and rent) to those controlling energy production.

- **Intellectual Property and Licensing:** When creative and design labor is automated, **intellectual property rights** become a primary tool for value capture. Consider entertainment media: if AI can generate films, music, and art on demand, the only thing preventing a flood of interchangeable content is IP law (copyrights, trademarks, etc.) and branding. Already, major content owners are moving to license their catalogs to AI model developers <sup>23</sup> or to restrict unauthorized use of their IP in AI outputs. In a fully AI world, a company like Disney can leverage its famous characters and franchises – even if *anyone* could technically produce a Star Wars-like film with AI, *only Disney's version is legal and authentic*. Thus, corporate control of popular IP will let those firms dominate profits in content, while generic AI-made content may struggle to monetize if it infringes or if audiences gravitate to known brands. Similarly in technology and manufacturing, patents and trade secrets on designs or processes will be goldmines. If a breakthrough AI-developed drug can be manufactured by automated labs, the *formula patent* or regulatory approval becomes the source of monopoly profit, not the manufacturing labor. Patent holders could license their designs to countless robot-run factories worldwide and collect royalties with minimal marginal cost. We might also see **new IP-like regimes** – for instance, databases of training data might get sui generis protection, or companies might secure exclusive rights to certain AI algorithms. On the flip side, if AI-generated outputs aren't granted IP protection (as current law leans <sup>12</sup> ), some products could become commodities. But big firms will find ways to maintain proprietary edges: e.g. by embedding subtle human creative input to claim copyright, by trademarking AI-generated marks, or by lobbying for expanded IP laws to cover AI creations. *Licensing agreements* will proliferate – whether it's licensing software (AI models as a service), licensing content to AIs, or licensing the AI's output. All of these create rent streams for the licensors (often incumbent firms or AI platform providers).
- **Regulatory Permissions and Rule-Making:** Regulation can act as a gatekeeper that limits competition. In an economy of abundant AI, **permission to operate in certain sectors** may become the hardest barrier. For example, fully autonomous vehicles have been technically ready to drive in many areas, but their rollout is gated by regulators granting permits and approving safety. Companies that navigate this regulatory maze quickly (often well-resourced incumbents) gain a lead. Additionally, large firms often influence regulations to **raise barriers to entry** – this is a classic phenomenon where complex compliance requirements favor those with armies of lawyers and lobbyists <sup>15</sup> <sup>17</sup> . In finance, for instance, an AI-run hedge fund or bank still needs licensing from authorities; big banks can fulfill capital and reporting requirements that a small AI startup cannot. If governments impose *AI certification exams*, *audit standards*, or *reporting mandates* (say, any AI model above X computations must undergo oversight <sup>27</sup> ), the result may be fewer, more dominant players who can afford compliance. While such regulation is often well-intentioned for safety, it incidentally **protects incumbents**. Moreover, governments themselves might become chokepoints: for instance, export controls on advanced AI chips (already seen in US-China trade policy) mean only certain nations or companies can access top hardware <sup>26</sup> . Regulatory permits in sectors like healthcare (e.g. FDA approval of an AI diagnostic device) or aviation (drone delivery flight corridors) will decide who can monetize AI innovations. We may even see *franchising via regulation* – governments could auction limited licenses for AI services (imagine only a set number of AI healthcare providers allowed nationally for quality control), thereby creating scarcity value for license holders. In summary, **the labyrinth of laws and approvals becomes a competitive battleground**, and those who secure compliance and influence law will hold advantageous positions. This chokepoint is somewhat under human control (lawmakers), but it translates into economic power for firms that align with or shape the regulatory regime.

• **Scarce Physical Assets and Resources:** Lastly, the physical world still imposes scarcity. Land, raw materials, and certain key infrastructure cannot be infinitely replicated by AI. If robots make manufacturing cheap, the limiting factor might be the **availability of raw materials** (minerals, rare earths, water, etc.) and the *ownership of land or facilities*. For example, fully automated agriculture could produce food very cheaply – but fertile land is finite, so farmland owners could capture much of the value of increased agricultural productivity (through land rents or high land prices). In mining and heavy industry, having rights to mineral deposits or oil fields becomes even more crucial if extraction is automated (no labor strikes, constant operation). The owners of these natural resources might reap outsized gains since robots simply enable tapping the resource faster and cheaper. We might see **vertical integration** where AI companies acquire resource producers to secure their supply (e.g. a battery factory AI buying lithium mines). Additionally, infrastructure like ports, railways, warehouses, telecom spectrum, etc., remains limited. A logistics company with automated trucks still needs warehouses near cities and access to highways; if those facilities are controlled by a few players or require permits, it's a chokepoint. In telecommunications, even if AI automates network management, the radio spectrum and fiber networks are finite and regulated – owners of those networks (or government license holders) maintain pricing power. Essentially, *any real-world bottleneck that automation cannot eliminate becomes a source of rent*. In a sense, economics might revert to a *landlord model*: those who own the non-replicable pieces of the world (land, resources, networks) can charge others for use, since labor is no longer the bottleneck cost.

In combination, these chokepoints suggest that while the cost of producing many goods and services will plummet (thanks to AI/robot efficiency), the distribution of gains will hinge on strategic control points. Wealth may concentrate even more in the hands of those who control capital, IP, data, compute, and critical assets, unless countervailing policies (antitrust, wealth redistribution, open-source movements) intervene. For consumers, some products might become extremely cheap and plentiful (pure digital goods, for example), but others might remain expensive if they are tied to scarce inputs or monopolized channels.

## V. Sectoral Transformations Around New Chokepoints

The impact of abundant AI labor will not be uniform across industries. Each sector will restructure around whichever chokepoints are most relevant to its value chain. Let's examine a few major sectors – **Entertainment, Finance, Logistics, and Heavy Industry** – to see how fully autonomous operations might reshape them:

### A. Entertainment and Media

The entertainment industry offers a vivid example of creative labor being disrupted by AI. In a future where AI models can write scripts, compose music, generate realistic films or games, the traditional studio staffed by writers, animators, directors, and crew could be largely replaced by *content generation engines*. We are already seeing the beginnings of this: in 2023, Hollywood writers and actors went on strike largely to secure protections against AI encroachment on their jobs <sup>28</sup>. They feared that **“unchecked AI could dramatically reshape Hollywood and undermine their roles, pitting artists against robots in a battle over human creativity.”** Their new contracts put guardrails – for instance, studios agreed not to use AI to generate scripts or as source material without proper compensation <sup>29</sup>. These guardrails, however, are temporary buffers. As AI gets better and if economic pressure mounts, studios (or new AI-native media companies) will push to use AI more aggressively to cut costs and increase content output.



**Structure of Entertainment Firms:** We can imagine an entertainment company in 2035 that has almost no creative staff. Instead, it has: a library of IP (characters, franchises, past designs), a powerful generative AI system that can create new movies/shows/music on demand, and a distribution platform (streaming service or VR world) to deliver content. The *core human roles* might be a few brand managers or editors who decide high-level direction (“Make the next superhero movie have these themes and meet PG-13 standards”) and perhaps a legal team to handle IP and rights. Everything else – scriptwriting, storyboarding, special effects, even casting virtual actors – could be handled by AI. The corporation thus looks more like a tech company than a traditional studio: its key assets are its IP rights and its algorithms, not a roster of creative talent.

**Chokepoints & Rents:** In entertainment, *creative labor* becomes cheap, but *creative IP* becomes precious. Companies that own beloved franchises, characters, or vast content libraries will leverage those to differentiate their AI-generated content. For example, Disney could automatically generate endless Star Wars or Marvel spin-offs using AI, and fans will watch because they recognize the brand – no human writers needed. Meanwhile, an independent creator might also use AI to make a space opera film of comparable quality, but they can’t legally use the Star Wars universe and may struggle to get audience attention in a saturated market. Thus, **ownership of popular IP and strong branding is a major chokehold**. We might see an arms race to acquire IP – classic literature, old film catalogs, game universes – to feed the AI and be allowed exclusive use of those settings and characters. Also, *distribution channels* remain important. Platforms like Netflix or YouTube (or future VR entertainment hubs) have captive audiences; even if content is cheap to make, getting eyeballs on it may require being on the dominant platform. Those platforms, in turn, might produce their own algorithmic content and could down-rank independent productions. In summary, entertainment firms will likely restructure into **lean IP-holding content factories**: minimal employees, heavy use of AI, and focus on monetizing story franchises, subscriber bases, and licensing deals.

The economics might shift to a *hits and subscription* model even more. If AI can churn out unlimited content, the value of any single piece drops – so companies will either aim to produce the *absolute top hits* (using marketing analytics and AI to craft content that maximizes engagement) or to lock consumers into ecosystems (e.g. a subscription where your personal AI curates a continuous stream of content for you, all proprietary to the platform). Human stars and influencers might still matter – genuine human artistry could become a *luxury good* in a sense – but even likenesses can be duplicated by AI with permission. We might see actors licensing digital replicas of themselves to studios, so the actor’s *image becomes IP* that generates income while the person doesn’t actually perform. Indeed, part of the SAG-AFTRA strike was about protection against use of actors’ digital likenesses. Once resolved, it’s conceivable that future contracts will allow studios to create new films starring an AI-generated version of an actor (who maybe records a minimal performance or just gives rights), enabling the actor to “appear” in many projects simultaneously. This again centralizes value on *those with rights to recognizable faces or characters*.

**Creative Process and Governance:** Traditional corporate governance in media involved greenlighting decisions, budgets, talent contracts, etc., managed by executives. In an AI-run studio, greenlighting might be as simple as running simulations: the AI could predict audience reactions and revenue for a proposed film, and automatically commence production if projections are favorable, adjusting content in real-time based on test audience AIs. The “executive” function becomes one of setting constraints (like compliance with censorship rules or brand guidelines) and letting the AI optimize. The board of such a company might mostly concern itself with public relations and regulatory issues (ensuring the content doesn’t trigger legal problems) rather than choosing creative direction.

**Risks and Opportunities:** Entertainment may flood with content – a blessing for consumers in choice, but also a risk of overwhelming noise. Societally, we might value *curation and authenticity* more when AI content is ubiquitous. This could give a niche role to humans – e.g. a “human-curated” film festival might become prestigious compared to algorithm-generated mass content. Companies could capitalize on this by branding some content as “human-made” artisanal content (ironically, a reversal where human labor is the special ingredient). However, the mainstream economics will favor volume and personalization: *tailoring content streams to each user via AI*, something only big platforms with data can do effectively. In conclusion, the entertainment sector will likely consolidate around major IP owners and distribution platforms, with drastically fewer creative employees. Profit margins might initially spike due to labor cost savings, but could later thin out if content becomes commoditized and competitive – unless kept high by the monopolistic control of IP and platforms.

## B. Finance and Banking

The finance sector is already highly digitized and algorithmic, making it ripe for full AI integration. Imagine a bank or investment firm where **AI handles all core functions**: algorithmic trading bots manage the investment portfolio, AI underwriters evaluate loan applications, robo-advisors handle customer wealth management, AI compliance systems monitor for fraud and regulatory issues, and even strategic decisions (like entering a new market or adjusting risk exposure) are made by analyzing massive data through AI models. Many of these elements exist today in parts – high-frequency trading is automated, big banks use AI for credit scoring and fraud detection, and chatbot advisors serve retail customers. The future corporation in finance could be a largely *server-based entity* with only a nominal human skeleton crew (perhaps just to satisfy regulators or maintain client relationships with large institutional customers who still like some human contact).

**Restructuring and Automation:** A fully automated financial firm might not look dramatically different outwardly (it will still offer loans, investments, insurance, etc.), but internally the human headcount could shrink by orders of magnitude. For example, a hedge fund might run on a proprietary AI that ingests all global market data and executes trades, with no human analysts or traders on staff. The “firm” could be just the AI, the computing infrastructure, and the capital it’s investing, overseen by a couple of technicians. In banking, retail operations could be entirely ATM- and app-based; physical branches may vanish or be staffed by a single IT person while AI systems handle transactions and customer inquiries. **Corporate structure** here might retain more human facade due to trust – customers might demand some accountability. But technically, even customer service can be AI avatars.

**Chokepoints & Economic Dynamics:** In finance, *capital* is obviously the key resource – an AI hedge fund is useless without a large fund to invest. So those with capital (big institutions, wealthy individuals) still hold the power. However, if AI makes generating returns easier or more equal, one might expect competition to drive down returns (the efficient-market hypothesis on steroids: if every firm has super-intelligent traders, excess profits from trading might vanish). Where, then, can margins persist? Possibly **scale and network**. For instance, large banks might use AI to achieve such efficiency that they can offer better rates and still profit, squeezing out smaller banks. We might see *even more consolidation* – if one AI-driven bank can serve 100 million customers with flawless efficiency, why have thousands of regional banks? It could lead to mega-banks (unless antitrust intervenes). Another chokehold is **data**: financial institutions with decades of customer transactional data can feed their AI for better predictive models (like anticipating who will default, or detecting market micro-structure patterns). New entrants without historical data might be at a disadvantage or have to buy data. Access to *real-time market data* is also crucial – exchanges might sell

premium data feeds at high prices, benefiting incumbents who can afford them and giving them a slight speed edge.

**Regulation in Finance:** Finance is heavily regulated to ensure stability and consumer protection. In a fully AI-run scenario, regulators will likely enforce strong oversight. *Regulatory permission* becomes a moat: only those entities that prove their AI is safe and compliant may be licensed as, say, a deposit-taking bank or a broker-dealer. The cost of such compliance (validating algorithms, continuous auditing, cybersecurity) is substantial, favoring big players. There's also systemic risk concerns – if many firms use similar AI strategies, could they all crash the market simultaneously? Regulators might require diversity of algorithms or impose brakes (like circuit-breakers). So ironically, even if tech would allow infinite new AI funds to pop up, regulators might *limit their number or scope* to prevent instability. That means **licenses or charters in finance become valuable limited commodities**, and existing big banks/insurance companies likely keep them, whereas new purely AI entrants struggle to get approved.

**Margin Opportunities:** One area where margins could persist is in **market-making and infrastructure**. Exchanges, clearinghouses, and payment networks could remain highly profitable because *everyone* (human or AI-run firms) must use them to transact. If those infrastructures are not fully commoditized, their owners can charge fees. For example, if trading is dominated by AIs, the exchanges might increase fees for data and execution because the AIs have no alternative (overseen to avoid overt exploitation, but some rent likely remains). Another is **private markets and proprietary deal access**: AI can analyze public markets easily, but having relationships or rights to certain deals (say a stake in a private company or a real asset investment) could yield returns that generic AI trading cannot access. Thus, large financial institutions might pivot more to *origination of deals and products* (like structuring complex derivatives or financing projects) where having a network and trust (and regulatory OK) matters. They'll use AI to optimize those deals, but the advantage is in being at the table, which newcomers can't easily replicate with just an algorithm.

**Employment and Governance:** For corporate governance, finance firms might retain human boards longer simply due to regulatory expectation (e.g., regulators might mandate that a human risk officer and board sign off on the AI's models to take responsibility). But those humans might rely heavily on AI analytics in making any decisions. Shareholders will still own the companies, but they'll likely demand new forms of assurance, such as *AI audit reports* and *contingency plans* if the AI malfunctions (flash crashes etc.). The fiduciary duty of financial advisors could be codified into AI – maybe we'll see legally “fiduciary AIs” that manage client money with a duty of loyalty programmed in. Already, some jurisdictions impose a fiduciary standard on robo-advisors as if they were human advisors. This will continue.

In sum, the finance sector may become **hyper-automated but still oligopolistic**. The big get bigger, leveraging AI at scale, while niche players survive by specializing or catering to specific needs (perhaps combining human insight with AI in unique ways for clients who value the human touch). Costs for routine transactions could plummet (benefiting consumers with cheaper banking and investment services), but profits for leading firms might remain high through sheer volume and control of the financial infrastructure. If cognition is commoditized, the winners in finance are those with *most capital, best data, regulatory clearance, and client trust* – all factors that current incumbents possess in abundance.

## C. Logistics and Supply Chain

Logistics – the movement of goods – stands to be dramatically streamlined by robotics and AI. Consider a supply chain in which **warehouses are fully robotic**, trucks drive themselves, cargo ships sail autonomously, and delivery drones or robots handle last-mile drops. The vision of the “lights-out” warehouse is already being pursued by companies like Amazon, where robots fetch and sort items with minimal human intervention. Self-driving trucks and robotaxi services are operational in pilot programs (e.g. Waymo and Cruise vehicles navigating city streets without drivers, thousands of rides per week) <sup>30</sup>. Once perfected and scaled, a logistics company could conceivably operate an end-to-end goods delivery network 24/7 with only a handful of human supervisors remotely monitoring the systems.

**Industry Restructuring:** Today’s logistics involves many players – trucking companies, shipping lines, 3PL (third-party logistics) providers, warehouse operators, etc. Automation could lead to **vertical integration** or at least the dominance of a few platform players. For instance, a giant retailer (like Amazon or Alibaba) might own its automated warehouses, a fleet of autonomous trucks and cargo drones, and use AI to coordinate the entire supply chain from factory to consumer. If their system is highly efficient and handles massive volume, it’s hard for smaller competitors to match the speed and cost. We might get a scenario where global logistics is run by a few large AI-powered networks (like a “Cisco of logistics” providing backbone services, or the retail giants themselves). Conversely, one could imagine decentralized networks if technology becomes cheap – e.g. small businesses using open-source AI and renting robot vehicles to ship goods autonomously in a more distributed fashion. However, even in that scenario, there will be centralized chokepoints like key transport hubs or digital platforms matching cargo to vehicles.

**Chokepoints & Key Assets:** In logistics, *physical infrastructure and regulatory permissions* are paramount. Autonomous vehicles need permission to operate on roads or airspace. Early approvals have been city-by-city; companies that win trust of regulators (through safety records and lobbying) have a head start. *Route data* and optimization algorithms are also crucial – companies with years of delivery data (like UPS or FedEx optimizing routes) have an advantage training AI for route planning. But perhaps the biggest asset is **infrastructure**: warehouses in strategic locations, access to ports, proprietary fulfillment centers near customers. These require capital and sometimes government cooperation (zoning, etc.). Big retail/logistics firms have been investing in these heavily. In an AI future, these become even more of a moat because without human labor needs, one giant distribution center can serve an entire region if placed optimally – leaving little room for multiple competing centers. So, **location and network effects** (the more extensive your delivery network, the faster and cheaper you can deliver, attracting more business, which further funds expansion) will likely drive consolidation.

Another chokepoint is **system integration**: coordinating thousands of moving pieces (vehicles, inventory levels, weather data, etc.) is complex, and a firm with a superior AI operations platform can outcompete. That platform itself might become a product: for example, a tech company could run a *logistics AI service* that other companies plug into (somewhat like how Amazon Web Services hosts many companies’ computing, perhaps “AWS for logistics” could route their shipments via its AI). The platform owner would then skim profit from all logistic transactions.

**Efficiency and Margins:** As labor costs drop out, the cost per delivery plummets. We could see extremely cheap shipping – maybe even essentially free next-day delivery as a standard, because the marginal cost of a drone drop is cents of electricity. However, initial capital and maintenance costs (and energy costs) remain. If competition is vibrant, those savings pass to consumers. If a few companies dominate, they might keep

prices higher to earn profit. Likely, given the physical nature of logistics, *economies of scale* will push toward a few large players, but some competition between them will keep rates reasonable. Margins then depend on how efficiently each can run their network relative to others. If one company's AI optimizes 5% better, at scale that yields big profit unless they undercut on price to gain market share.

**Employment and Governance:** The logistics sector currently employs millions of drivers, warehouse pickers, and so on. In the fully automated scenario, those jobs disappear, replaced by maintenance techs, fleet managers, and AI supervisors (far fewer roles). Corporations in this space might proactively rebrand from being “logistics employers” to “logistics platforms.” They might face regulatory or public pressures to maintain safety – e.g., requiring a human call center for emergency situations when an autonomous truck encounters something unexpected. Governance-wise, the companies will focus on *operational reliability*: boards will scrutinize metrics like system uptime, accident rates, delivery times, because any glitch in an AI logistics network could be costly (imagine a software bug grounding an entire fleet). Shareholders will demand continuous innovation to keep these AIs at peak efficiency. Perhaps partnerships will form – e.g., logistics firms partnering with energy companies (to secure power for electric vehicles) or with municipalities (smart city infrastructure to assist autonomy).

**Customer Impact:** For consumers and businesses, logistics automation means **faster and more predictable supply chains**. Inventory management may shift – if you can deliver anything in an hour with robots, you don’t need local stock, which reduces inventory costs (benefiting manufacturers and retailers). This could further entrench big players like Amazon who already operate on thin inventory via quick fulfillment. Small retailers might survive by tapping into the dominant networks (at the mercy of their pricing). Alternatively, local 3D printing and manufacturing (also automated) could shorten supply chains and somewhat reduce the volume of long-distance shipping, but that too would likely be run by companies integrated with logistic networks.

Overall, logistics in a fully automated world likely becomes **a utility-like sector dominated by a few AI-driven networks**. Those who control the major routes, warehouses, and digital platforms for moving goods will extract consistent rents, while the cost per delivery for the rest of society falls. The corporate form here might be mega-corporations controlling entire end-to-end pipelines, or a tight oligopoly of platform providers coordinating with each other (like airlines alliances, but even more consolidated). The focus for these corporations will be on *managing huge fleets of machines* and negotiating with regulators and local governments – essentially a high-tech infrastructure management game, with AI as the dispatcher.

## D. Heavy Industry and Manufacturing

Heavy industries – such as manufacturing, mining, construction, and energy production – have been marching toward automation for decades, from robotic assembly lines to self-operating mine trucks. In the scenario of AI/robotics “solved,” we reach the pinnacle of what’s often called “**lights-out manufacturing**” or dark factories, meaning factories that operate with no human workers on-site. This is not science fiction: for example, the Japanese robotics company FANUC has famously run a factory since 2001 where robots build other robots, unsupervised for weeks at a time, producing 50 units per day <sup>31</sup>. The concept has expanded to micro-molding, CNC machining cells, and other processes where humans only set up the machines and then the automated line runs continuously <sup>32</sup>.

In a fully realized vision, a heavy industry corporation might own a collection of such lights-out facilities. **Mines** could have automated drills and haul trucks (Rio Tinto has deployed autonomous haul trucks that

operate more efficiently than human-driven ones <sup>33</sup> ). **Factories** could run 24/7 with robotic arms, 3D printers, and AI QC (quality control) systems inspecting products via computer vision. **Construction** might see autonomous bulldozers, drones surveying sites, and robot bricklayers or 3D concrete printers erecting buildings with minimal human crew. The upshot is vastly increased productivity and output capacity – but concentrated in the hands of whoever owns these automated systems.

**Structural Changes:** Manufacturing companies historically required large human workforces, leading to considerations of labor relations, location (proximity to labor or low-cost labor markets), and training. In an AI-dominant era, those considerations give way to others: access to raw materials and energy, the regulatory environment (especially regarding safety and environmental rules), and capital for high-tech equipment. We may see **production geographies shift**. For instance, offshoring to low-wage countries could reverse; factories might return to consumer markets (like the U.S. or Europe) because cheap labor is no longer a factor, while proximity to customers for faster delivery and stable governance is attractive. However, access to raw materials might keep some production near resource sites (an automated textile factory might still be near cotton fields or a port for imported inputs). *National policies* could also shape this: countries will compete for hosting these advanced automated industries by offering tax breaks or infrastructure, since jobs are not the carrot, perhaps the promise of tech leadership or supply chain security is.

**Chokepoints & Ownership:** In heavy industry, **ownership of production means** (the factories, plants, and robots themselves) becomes central. If previously a factory's value was partly in its skilled workforce, now it's almost entirely in its machinery and software. These are capital-intensive to set up – indeed, a fully automated plant is only cost-effective with very high volume and careful process engineering <sup>34</sup> . This suggests fewer but larger plants for many products (to maximize economies of scale). Companies that can *afford and manage* such large-scale automation will dominate their industries. Smaller factories might not compete unless they find niche bespoke products or leverage flexible automation for high-mix low-volume production (AI can help with that too, but complexity favors those who master it).

Another chokepoint is **intellectual property and know-how** in manufacturing processes. The combination of AI and robotics to make a complex product (like advanced semiconductors or pharmaceuticals) is itself a competitive advantage that can be protected via patents or trade secrets. If one company's AI factory runs at 10% lower cost due to better algorithms or proprietary machine designs, that company can undercut others or enjoy higher margins. There might also be *network effects* in manufacturing – for example, a platform that connects automated factories to designers globally (like a cloud manufacturing service) might become the go-to, thereby capturing a slice of all production.

**Resource Control:** As mentioned, *raw materials* are vital. Mining could be revolutionized by robots, but you still need to own mineral rights. If anything, automation might make previously marginal mines viable (no need to attract workers to remote or dangerous areas). Companies could swarm to secure mining rights, leading to an enclosure of resource sources. Commodities might drop in price with easier extraction, but if demand remains high, the resource owners still profit. The same logic applies to energy – automated drilling or renewable energy farms still need land and resource rights. With high automation, energy supply might increase (cheaper solar/wind farms maintained by robots), potentially lowering energy costs, but the control of energy assets is strategic (we see countries and conglomerates already treating battery supply chains or rare earth minerals as strategic assets).

**Sector Integration:** Heavy industry corporations might integrate vertically to secure their chokepoints. For example, an electric vehicle company might not only automate car factories but also automate battery production and lithium mining under one corporate umbrella, ensuring it controls the full stack from resource to finished product. If each stage is AI-run, coordinating them under one roof can eliminate supplier margins and bottlenecks. This could challenge the traditional model of specialized firms in a supply chain – why buy parts from a supplier if your robots can make the parts in-house just as cheaply? The counterpoint is flexibility: maybe a central AI platform could coordinate a *virtual vertically integrated supply chain* among multiple companies dynamically. But such coordination likely falls to whoever has the greatest reach – possibly the same large corporation or a platform provider.

**Labor and Skills:** While direct labor jobs shrink, the need for *highly skilled technical professionals* initially rises – people to design the automation, program the AIs, maintain complex equipment, etc. Over time, AI might handle much of the *design* as well (AIs designing better robots and factories). Maintenance could be done by other robots (self-repairing machines, or at least robots that swap out parts for each other). Still, heavy industry often has unpredictable elements (geology in mining, weather in construction) that require some human problem-solving in the loop for the foreseeable future. Corporations will likely keep or contract a small workforce of engineers who oversee multiple automated sites. We may see **remote operations centers** where a handful of experts monitor dozens of mines or plants via AI dashboards, stepping in only as needed. Already, mining companies have remote operation centers for some truck fleets.

**Governance and Externalities:** Heavy industries have significant external impacts – environmental pollution, resource depletion, safety hazards. If not managed responsibly, fully automated operations could, for instance, mine 24/7 and cause faster environmental damage than human-paced operations (since no need to stop for shifts). This raises the importance of regulation and corporate governance focusing on *sustainability and safety*. Boards of these companies will need to ensure compliance with environmental laws and possibly ethical guidelines – e.g., preventing an AI from maximizing output at the expense of an ecological disaster. Liability for accidents in heavy industry (like a chemical plant explosion) will still lie with the company; avoiding such accidents when humans aren't present might mean heavy reliance on sensor systems and AI predictive maintenance. There's potential for **safer operations** (robots don't get tired or lax on safety rules, and no workers put in harm's way), but also for new failure modes if AI misjudges a situation. Governance will thus involve a strong *engineering risk management* perspective on the board.

**Competition and Margins:** Heavy industry could go one of two ways: either ultra-efficient commodity production with thin margins (if many can do it) or winner-takes-most situations where the first movers or best tech yield significantly lower costs and can outlive others. Taking manufacturing as example, if multiple firms adopt lights-out factories, they might compete on price, benefiting consumers with cheaper goods – until perhaps only the most efficient survive in each product category. In mining, if commodities become oversupplied due to easier extraction, prices fall, and only the lowest-cost producers (with the best tech or richest ore grade) stay profitable. Thus, *even with automation, basic economic forces apply*: competition drives prices toward cost. The difference is cost now largely depends on technology and asset quality rather than labor efficiency. So, margin will accrue to those with **proprietary tech (cost advantage)** or controlling a **higher value product niche** or a resource monopoly. For instance, a company with a patent on a superior automated refining process could produce higher-purity metal at lower cost than others and enjoy a premium until others catch up or license the tech (collecting royalties).

**Example – A Future Steel Company:** To illustrate, envision a steel manufacturing corporation circa 2040. It owns iron ore mines in Australia operated by autonomous drills and haul trucks, a fleet of AI-run cargo ships that transport ore to its steel plants (the ships navigate and self-maintain), and “dark” steel mills where raw ore is processed by AI-managed furnaces and rolling lines with robotic handling. The entire flow from mine to finished steel coils is algorithmically scheduled to minimize energy use and downtime. This company employs maybe 100 human experts worldwide to monitor systems and handle exceptional cases, versus 50,000 workers in a comparable 20th-century steel conglomerate. Its competitive advantage lies in owning high-grade iron ore reserves and having proprietary AI process controls that make its steel stronger and cheaper to produce. It has essentially *no labor costs*, so its cost per ton is extremely low – yet if there are only a few such players globally, they keep prices just low enough to deter new entrants but high enough to make substantial profit (with profits going to shareholders and reinvestment in even more efficient tech). If steel demand is high, the company thrives; if oversupply occurs, it might idle some plants (the cost of idling is lower when you don’t have to pay salaries, you just turn off robots, though you still have capital depreciation). This company’s main concerns: securing enough electricity (maybe it builds its own solar farms), maintaining good relations with governments (to keep mining rights and navigate any protectionist trade policies), and staying ahead in the tech race (to prevent a competitor from underpricing them with a better AI or process).

In conclusion, heavy industry will likely be characterized by **high capital intensity, low direct labor, and a premium on technology and resource ownership**. The number of players in each field might shrink (because you don’t need as many factories if each is hyper-productive), leading to oligopolies in some markets. Alternatively, if tech is widespread (e.g., many companies license the same AI factory tech), it could remain competitive but with all firms having similar structures. Either way, the *purpose of corporations here remains the production of physical goods*, but the way they achieve it – through automated assets – alters the scale and scope of each corporation. They become more like *machine networks* than human organizations.

---

Across all these sectors, a clear pattern emerges: **the sources of value move away from human labor and toward the control of scarce inputs and the ownership of advanced systems**. Corporations, in adapting to this, may become fewer, larger, and more entwined with technology providers and regulators. The economy might produce far more output with far fewer people directly involved, challenging how we think about everything from competition to employment to wealth distribution. The next section concludes with overarching implications and theoretical insights from legal and economic research on this radical transformation.

## **Conclusion: Corporations Beyond Labor – Efficiency, Concentration, and New Responsibilities**

In a world where AI and robotics can do essentially all work, corporations remain indispensable – but their **purpose and design shift fundamentally**. They no longer primarily serve as structures to coordinate human workers, but rather as *owners and orchestrators of automation*. Key aspects of corporate law (entity personhood, limited liability, ownership shares) persist as the scaffolding allowing autonomous operations to interact in markets. However, many traditional governance features (human boards, executive teams, employee hierarchies) may be pared down or reimagined to accommodate algorithmic management and decision-making.



From a **legal perspective**, this scenario pushes the boundaries of our concepts of personhood and responsibility. The law will need to ensure that human accountability and ethical norms don't vanish when humans step out of direct control. This might involve new regulations – such as requiring corporations to disclose AI decision criteria, carry insurance for algorithmic harms, or even to embed certain fiduciary-like constraints into their AIs. Corporate personhood might expand to cover autonomous agents, or conversely, the corporate entity will act as the legal “person” shielding and channeling the acts of non-human agents within. Current experiments with “DAOs” (decentralized autonomous organizations) and memberless LLCs indicate that our legal system can flex to allow non-traditional corporate forms <sup>2</sup> <sup>3</sup>, though not without resistance or need for clarification. Crucially, *property rights and limited liability* will continue to provide the incentive for innovation – entrepreneurs will deploy AI businesses knowing their personal assets aren't at stake beyond their investment, and that the AI entity can own property and sue/be sued like any person.

From an **economic standpoint**, the commoditization of labor leads to a search for new scarcities. We identified capital, data, IP, compute, regulatory access, and physical assets as the new strategic assets. Economic theory suggests that rents (excess profits) accrue to scarcities protected by barriers to entry. In the fully automated economy, **market power may concentrate intensely** in the hands of those who control these bottlenecks. Early evidence of this is visible: large tech-centric firms are becoming “superstar” companies partly because AI investments allow them to scale faster and dominate markets <sup>20</sup>. With AI as an equalizer of skill, advantages of scale and incumbency (like having more data or capital) grow more important, potentially leading to natural monopolies or oligopolies in various sectors.

This raises policy questions: Will antitrust law need to intervene more aggressively if, say, only a couple of companies effectively run global logistics or finance via AI? Or will new competitors constantly emerge since AI lowers entry barriers in some respects (an AI can power a startup's operations from day one)? There's an argument that *open-source AI and democratized technology* could allow smaller players to compete – for instance, if computing power becomes cheap and models widely available, maybe thousands of micro-factories and small AI firms thrive, preventing total concentration. The reality might be a mix: core infrastructure and platforms consolidate, while niche or creative endeavors proliferate on top of them. The outcome depends on deliberate choices in technology sharing, regulation, and possibly the emergence of *cooperative models* (like public or community-owned AI utilities) as a counterweight to corporate concentration.

Finally, the **social role of corporations** may need redefinition. If firms no longer provide mass employment, governments and society might expect them to contribute in other ways – perhaps through higher corporate taxes to fund social safety nets (like a universal basic income, as the EU report hinted <sup>19</sup>), or through direct provision of public goods. Corporations themselves might adopt missions beyond profit if profit becomes too easy (imagine extremely profitable AI companies turning to ambitious projects like space colonization or climate engineering simply because they can). Alternatively, if unchecked, corporations might single-mindedly pursue profit with even less human empathy than before, since AIs won't have an innate sense of social responsibility. This could exacerbate issues like environmental degradation or exploitative consumer practices unless curbed by regulation or corporate ethics frameworks. Paradoxically, removing human workers could free companies to be *more humane* in some ways (no workers to mistreat), but it could also remove an important stakeholder that used to internally advocate for fair practices (employees often push companies toward community engagement, sustainability, etc.). The mantle of guiding corporate behavior might shift more to governments, consumers, and investors – for example, socially conscious investment funds could program their AI trading algorithms to favor companies with positive impact, indirectly influencing corporate priorities.

In summary, the era of effectively solved AI and robotics heralds a **new corporate paradigm**: firms become *hubs of capital and intelligence* rather than labor, and the critical governance questions revolve around controlling that intelligence and its impacts. The structure and purpose of corporations will adapt – preserving core legal benefits like asset pooling, risk shielding, and continuity, while shedding or modifying the human management apparatus that was once essential. The law will face the challenge of *keeping humans “in the loop” of accountability* without stifling the efficiency gains of automation. Economically, society stands to gain immense productivity and wealth, but the distribution of those gains will depend on how we manage the emerging chokepoints. Will we see widespread prosperity with cheap goods and services, or a techno-feudal concentration of wealth? That outcome may rest on policy decisions made as this transition unfolds.

One thing is certain: **corporations will not disappear** simply because human workers do. Instead, they will evolve into perhaps the most powerful tools ever seen – *autonomous wealth-generating entities* operating at speeds and scales no human organization could match. Harnessing these entities for broad benefit, and ensuring they remain our servants rather than our unwitting overlords, is the next great project for corporate law, economics, and society at large. The groundwork is being laid now in legal theories of autonomous entities and economic analyses of AI’s impact <sup>1</sup> <sup>20</sup> , but the real test will come as theory becomes reality in boardrooms (and server rooms) around the world.

#### Sources:

- Shawn Bayern, “Autonomous Legal Entities are Already Possible Under American Law” – Oxford Business Law Blog (2019) <sup>1</sup> <sup>2</sup> <sup>3</sup> .
- Stanford Technology Law Review, “Entity Law for Autonomous Systems” (Bayern, 2017) – on corporate law flexibility and eliminating human boards <sup>4</sup> <sup>5</sup> <sup>6</sup> .
- European Parliament Legal Affairs Committee, report on Civil Law Rules on Robotics (2017) – proposing electronic personhood and related regulations <sup>8</sup> <sup>9</sup> <sup>35</sup> .
- U.S. Copyright Office, “Copyright and AI” Report (2025) – confirming that fully AI-generated works lack human authorship and are not copyrightable <sup>12</sup> .
- Berkeley Haas Research (Fedyk et al., 2020), *Investing in AI* – finding AI adoption concentrated in large firms and contributing to market concentration <sup>20</sup> <sup>22</sup> .
- AI Now Institute report (Vipra & West, 2023), *Compute and AI* – on the monopolization of computing power and its role in concentrating industry power <sup>24</sup> <sup>25</sup> .
- Fast Company (Melendez, 2024), “In the AI era, data is gold” – documenting the importance of data and deals to license data for AI model training <sup>23</sup> .
- R. Threlfall, DataScienceCentral (2023), “Regulation favours the incumbent” – discussing how compliance costs and licensing can entrench large companies in AI development <sup>16</sup> <sup>17</sup> .
- AMT Online (McGrew, 2021), “Lights-Out Manufacturing” – example of FANUC’s fully automated factory and the cost considerations of complete automation <sup>31</sup> <sup>34</sup> .
- The Guardian (Anguiano & Beckett, 2023), “How Hollywood writers triumphed over AI” – on the entertainment industry’s response to AI and the potential for AI to replace creative roles <sup>28</sup> <sup>29</sup> .

1 2 3 7 Autonomous Legal Entities are Already Possible Under American Law | Oxford Law Blogs

<https://blogs.law.ox.ac.uk/business-law-blog/blog/2019/11/autonomous-legal-entities-are-already-possible-under-american-law>

4 5 6 Microsoft Word - 19-1-4-bayern-final.docx

[https://law.stanford.edu/wp-content/uploads/2017/11/19-1-4-bayern-final\\_0.pdf](https://law.stanford.edu/wp-content/uploads/2017/11/19-1-4-bayern-final_0.pdf)

8 9 10 11 13 19 35 Give robots 'personhood' status, EU committee argues | Technology | The Guardian

<https://www.theguardian.com/technology/2017/jan/12/give-robots-personhood-status-eu-committee-argues>

12 Copyrightability of AI Outputs: U.S. Copyright Office Analyzes Human Authorship Requirement | Insights | Jones Day

<https://www.jonesday.com/en/insights/2025/02/copyrightability-of-ai-outputs-us-copyright-office-analyzes-human-authorship-requirement>

14 Employed Algorithms: A Labor Model of Corporate Liability for AI

<https://scholarship.law.duke.edu/dlj/vol72/iss4/2/>

15 16 17 27 Regulatory Capture: Why AI regulation favours the incumbents - DataScienceCentral.com

<https://www.datasciencecentral.com/regulatory-capture-why-ai-regulation-favours-the-incumbents/>

18 [PDF] Artificial Fiduciaries

<https://scholarlycommons.law.wlu.edu/cgi/viewcontent.cgi?article=4885&context=wlulr>

20 21 22 Investing in AI boosts firm growth—and increases market dominance, study finds - Haas News | Berkeley Haas

<https://newsroom.haas.berkeley.edu/research/investing-in-artificial-intelligence-ai-boosts-growth-and-increases-market-dominance/>

23 In the AI era, data is gold. And these companies are striking it rich - Fast Company

<https://www.fastcompany.com/91148997/data-is-gold-in-ai-era>

24 25 26 Computational Power and AI - AI Now Institute

<https://ainowinstitute.org/publications/compute-and-ai>

28 29 How Hollywood writers triumphed over AI – and why it matters | US writers' strike 2023 | The Guardian

<https://www.theguardian.com/culture/2023/oct/01/hollywood-writers-strike-artificial-intelligence>

30 The 2025 AI Index Report | Stanford HAI

<https://hai.stanford.edu/ai-index/2025-ai-index-report>

31 32 34 Lights-Out Manufacturing

<https://www.amtonline.org/article/lights-out-manufacturing>

33 Rio Tinto, BHP Reap the Benefits of Mining Automation - Mining Digital

<https://miningdigital.com/digital-transformation/what-are-the-benefits-of-automation-in-mining-operations>