ChatGPT

From S-Curves to New Frontiers: The Pace of Tech Adoption

The S-Curve of Technology Adoption – A Pattern Repeats

Technological revolutions often burst onto the scene amid hype and skepticism, yet their deep societal impact tends to emerge more gradually – and then all at once. Over decades of innovation, a consistent **S-curve pattern** of adoption has played out: an initial slow burn as a new technology proves itself, a steep climb as it hits mainstream relevance, and finally a plateau as saturation is reached [1]. From the rise of the personal computer and the internet in the late 20th century, to the more recent booms in smartphones, cloud computing, and electric vehicles, adoption curves have followed this familiar trajectory. Early on, only a niche group of innovators or enthusiasts use the new tool; then a tipping point is reached – often catalyzed by a breakthrough product or a convergence of factors – after which uptake accelerates dramatically. Eventually, as the market approaches saturation, growth tapers off. Each technology's timeline differs, but the broad arc of the S-curve has proven to be "one of the most frequent patterns in innovation" [2].

Understanding these historical S-curves is more than an academic exercise; it provides a lens to gauge how emerging innovations today might diffuse through society. For instance, the **smartphone revolution** was ignited when Apple's first iPhone in 2007 transformed a business executive's gadget into a mass-market necessity [3]. Over the next decade, smartphones went from a luxury for the few to an essential for billions, illustrating the classic S-curve from novelty to ubiquity. Likewise, enterprise technologies such as **data center virtualization** and **cloud computing** saw years of experimentation and hype before quietly reaching a "tipping point" where adoption surged – often at the moment when public enthusiasm had cooled and dedicated practitioners were refining the technology's real uses [4]. And with **electric vehicles (EVs)** and **industrial robots**, we have witnessed long gestation periods followed by recent exponential growth as costs fell, policies shifted, and the technologies proved their worth at scale. Each case offers empirical data on how quickly (or slowly) transformative change can happen once the conditions are right.

Today, we stand on the cusp of **new S-curves** in technologies that until very recently sounded like science fiction. Advanced generative AI – capable of performing complex cognitive tasks – is now being deployed as "copilot" assistants in everyday software. Humanoid robots, long a staple of futuristic fantasies, are taking their first tentative steps from research labs onto factory floors. The **second half of this report** will examine these nascent trends: how AI-based software agents integrated into tools like Office 365 or Google Workspace could achieve astonishing scale in record time, and how general-purpose robots like Tesla's Optimus or Figure's humanoids might ramp up production in the coming years. But first, to ground our understanding, we turn to the historical record – a tour of key technologies that reshaped our world, and the empirical ramp-up data that reveals how their adoption unfolded across regions from the United States and Europe to China, Japan, the UK and India.

Smartphones: From Zero to Billions in a Decade

In January 2007, Apple CEO Steve Jobs introduced the iPhone, declaring "today, Apple is going to reinvent the phone." At that time, smartphones – defined by internet connectivity and advanced operating systems – were a niche product. BlackBerry devices and Palm Treos had gained traction with professionals, but the **global mobile phone market** was still dominated by simpler feature phones. The iPhone's debut is often cited as a **tipping point for smartphones**, marking the moment the technology entered its rapid-growth phase [3]. Indeed, the subsequent adoption figures trace a textbook S-curve. **Within a decade of the iPhone's launch, smartphones went from a few percent of global phone sales to the majority**. By 2018, about **3 billion people were using a smartphone** – a stunning expansion from the technology's niche origins [3]. Fast forward to the mid-2020s, and an estimated **5.5 billion users** – roughly two-thirds of humanity – now carry these devices [3]. In classic S-curve fashion, annual smartphone sales have plateaued in recent years as most of the addressable market has been captured and many countries approach saturation [3].

This worldwide explosion in smartphone adoption did not occur uniformly. Different regions climbed the curve at different paces, reflecting local economics, demographics, and competitive dynamics. In the **United States**, smartphone penetration (the share of the population owning a smartphone) was only around 20% in 2010 [5]. But as Android handsets flooded the market and carriers rolled out 3G data networks, adoption surged. By 2011, roughly one-third of U.S. adults had a smartphone; by 2018, about **70%** did [5]. Today, over **91% of American adults own a smartphone**, a remarkable increase from just 35% in 2011 [6]. This trajectory – from one-third to nine-in-ten in the span of about a decade – illustrates how quickly the technology went mainstream once the value proposition (mobile internet, apps, and constant connectivity) became clear and affordable. A similar story played out in **Western Europe**. In the **United Kingdom**, for example, smartphone uptake accelerated in the early 2010s and reached roughly **79% of the population by 2020** [7]. Major EU economies like Germany and France saw comparable figures, around 75–80% penetration by 2020 [7], indicating that by the end of that decade most Europeans who wanted a smartphone had obtained one. Japan, which had pioneered advanced mobile phones earlier (with features like email and web access on flip phones in the 2000s), was initially slower to embrace smartphones but eventually joined the trend: by 2021 about **75% of Japan's population were using smartphones** [8], up from effectively zero a decade prior as consumers transitioned from feature phones to iOS and Android devices.

Crucially, the smartphone S-curve bent even more steeply in emerging markets once low-cost models became available. **China** offers a vivid case. In the late 2000s, China had hundreds of millions of mobile phone users but very few on smartphones. That changed rapidly with the proliferation of inexpensive Android phones and the expansion of 3G/4G networks. From virtually negligible levels in 2007, China's smartphone penetration shot up to **around 60% of its 1.4 billion population by 2021** [9]. By that year China had **865 million smartphone users** [9] – more than the entire population of Europe – and it continues to grow. Chinese manufacturers like Huawei, Xiaomi, and Oppo drove prices down and flooded the market, enabling urban and even rural consumers to leapfrog directly to smartphones. India's curve has been a bit more delayed but is following the same direction. In 2010, only a tiny fraction of India's 1.3+ billion people had smartphones; even by 2016, basic mobile phones still dominated. But with the rollout of ultra-cheap data plans (e.g. Jio's disruptive entry) and sub-$100 handsets, smartphone adoption in **India** rocketed upward in the late 2010s. By 2021 India had over **600 million smartphone users**, about **43.5% of the population** [9]. While that penetration rate lagged China's, it represented a huge jump from near-zero a decade earlier, and the absolute number of users in India is second only to China. Importantly, India's smartphone growth is far from finished – hundreds of millions of Indians have yet to upgrade from basic mobile phones, suggesting the upward climb of the S-curve is still ongoing there.

These numbers underscore the sheer speed of the smartphone revolution. After the early 2010s, adoption curves in country after country turned nearly vertical. For example, **U.S. smartphone ownership climbed from 35% of adults in 2011 to 91% in 2023** [6]. In **China**, the share of new phone sales that were smartphones went from almost nothing in 2009 to essentially 100% by the late 2010s as the country all but phased out non-smart devices. **Globally, the smartphone penetration rate rose from about 20% of people in 2010 to 72% by 2020** [10], with **6.7 billion** smartphone subscriptions estimated in 2024 [11]. By the end of the 2010s, annual smartphone shipments were in the billions of units, and markets like **Europe, China, and the U.S. became saturated**, with sales leveling off as most consumers already owned one [3]. The S-curve had reached its plateau in those advanced markets, even as late-adopting populations in places like rural India or sub-Saharan Africa continued climbing the curve. Indeed, smartphone adoption worldwide may be approaching a natural ceiling (perhaps in the ballpark of 80–90% of humanity), constrained only by the youngest and oldest cohorts and those in the most remote areas. The **smartphone S-curve, roughly 15 years in the making, stands as one of the fastest and most pervasive technology adoption waves in history**. It transformed how people communicate, work, and live – and it did so in roughly the span of a generation.

Cloud Computing and Virtualization: From Slow Start to Ubiquity

In contrast to flashy consumer gadgets like smartphones, **enterprise IT infrastructure** might seem a sedate arena. Yet the adoption of cloud computing – underpinned by server virtualization – has followed its own dramatic S-curve, one that was arguably **underappreciated in its early years**. When Amazon Web Services (AWS) launched its first cloud services in 2006, the idea of renting computing on demand was novel. For a long time, large corporations were hesitant to move mission-critical systems off-premises. As late as 2015 – almost a decade after the cloud's debut – only about **5% of enterprise IT workloads had migrated to the cloud** [12]. Initial adoption was slow and cautious, a classic "innovation plateau" phase where experimentation ran ahead of broad deployment. But behind the scenes, a crucial shift was taking place in data centers: **virtualization technology** was being rapidly implemented, setting the stage for cloud scalability.

Virtualization – the ability to run multiple virtual machines on one physical server – had existed in rudimentary form for decades, but it **broke out in mainstream use around 2009–2010**. The timing was no accident. The late-2000s global recession put pressure on IT budgets, and companies eagerly embraced virtualization to consolidate servers and cut costs. VMware's hypervisors and rival platforms matured to handle production workloads reliably. As a result, **data center virtualization surged from limited use to near-ubiquity over the 2010s**. By **2014, more than half of all server workloads were running in virtual machines**, according to Gartner analysts, and they projected that figure would reach 86% by 2016 [13]. In other words, within roughly five years, enterprise IT went from mostly physical servers to a majority virtualized environment. The implications were profound: virtualization not only improved efficiency but also made organizations cloud-ready. Once applications were decoupled from specific hardware, moving them to an off-premise cloud became far more feasible.

The **period from 2009 to 2020** can thus be seen as the virtualization S-curve, rising from niche to saturation. Industry surveys recount how initially only low-tier or non-critical applications were put on virtual servers; as confidence grew, even **business-critical workloads were virtualized** and by 2020 the practice was considered standard IT hygiene [14] [15]. Indeed, by the end of the decade, many enterprises had virtualized essentially everything that made economic sense, marking a saturation point for the technology's deployment. This quiet revolution within IT departments enabled the next phase: a **mass migration to the cloud**. As one venture analysis observed, **cloud uptake by large corporations was "very slow initially," but then accelerated rapidly** – roughly **nine years after AWS's launch came a turning point** [12]. In 2015 only 5% of workloads were cloud-based; by the early 2020s, approximately **60% of enterprise IT workloads are running in the cloud** [12]. This astounding leap within less than a decade reflects how quickly cloud services went from fringe to mainstream once the model proved its reliability and cost advantages. **Cloud adoption grew roughly twelvefold from mid-decade to 2023** [12], an emphatic steep portion of the S-curve.

Regionally, the cloud revolution has seen the **United States and Europe as early leaders**, with U.S. tech firms not only providing much of the cloud infrastructure (Amazon, Microsoft, Google) but also being among the first to adopt it for their own operations. European enterprises followed a similar trajectory, albeit with a bit more caution due to data sovereignty concerns and legacy investments. By the late 2010s, however, European adoption of cloud and virtualization was in full swing as well [16]. **Asia** presents a mixed picture: Japan's conservative corporate culture initially slowed cloud uptake, but that began to change in the 2010s. The standout is **China**, which built its own cloud giants (such as Alibaba Cloud and Tencent Cloud) and saw a rapid migration, especially among its burgeoning tech sector and internet companies. By the early 2020s, China had heavily invested in both public and private cloud capabilities, although lagging slightly behind the West in percentage of workloads migrated – partly due to later start and different industry mix. Meanwhile, **India** began embracing cloud services in the late 2010s, especially among startups and IT services companies, but many traditional industries are still in earlier phases of cloud adoption. Nonetheless, the global trend is clear: **we have passed the midpoint of the cloud S-curve**, and the COVID-19 pandemic only reinforced this by pushing even reluctant organizations to move operations online.

A vivid illustration of this trend is how **the hype cycle gave way to reality**. Cloud computing was over-hyped in its early years, then faced a period of sober reassessment (as companies grappled with security and cost concerns), but ultimately the **deployment caught up with the promises**, following the S-curve in converging with earlier high expectations [17]. Ironically, many enterprises truly committed to cloud modernization at a moment when it was no longer "fashionable" to tout – the noise had died down, and pragmatic solutions emerged. This pattern – **slow start, sudden acceleration** – is mirrored in the data. By 2023, **about 60% of enterprise workloads are cloud-based** [12], and major IT consultancies project that percentage will continue to rise toward the 80–90% range later in the 2020s as legacy systems are finally retired. The S-curve for cloud computing is thus still climbing but showing signs of approaching its upper asymptote in leading regions. One can foresee a time soon when essentially all organizations use cloud in some form (public, private, or hybrid) for practically all their computing needs – a plateau of near-total adoption.

Electric Vehicles: Accelerating Up the Curve

For much of the 20th century, the automobile industry was static in one respect: cars ran on internal combustion engines, burning gasoline or diesel. The idea of electric vehicles (EVs) has existed since the 1800s (early primitive electric carriages predate Ford's Model T), but modern EVs remained a curiosity until recently. In the 1990s and 2000s, a few limited models (like the GM EV1 or early Toyota Prius plug-in experiments) appeared but had negligible market impact. It wasn't until the **2010s** that EVs truly began their ascent – and now, in the 2020s, they are clearly in the steep climb of the adoption curve. **Global electric car sales** have grown exponentially in the past few years. In 2018, EVs (battery-electric plus plug-in hybrids) were only about **2% of new cars sold worldwide** [18]. By 2022, that share had jumped to 14%, and in 2023 it reached roughly **18% of global new car sales** [18]. In absolute terms, nearly **14 million** electric cars were sold in 2023, more than six times the number just five years earlier [19] [20]. This explosive growth indicates that the **EV market has entered the rapid growth phase of the S-curve**, moving beyond the early-adopter niche of affluent environmentalists and into the mass market.

Yet, as with other technologies, this global picture masks significant regional differences – essentially **multiple S-curves playing out at different speeds**. Nowhere has the EV transition been faster than in **China** and parts of **Europe**. China in particular is astonishing: as of 2023, **more than one in three new cars sold in China is electric** [21]. The Chinese government's aggressive incentives and domestic automakers' investments (from BYD to SAIC) have made China the largest and most dynamic EV market on earth. **China accounted for about 60% of global electric car sales in 2023** [22], and its EV sales grew 35% year-on-year even after subsidies ended, indicating an entrenched momentum [23]. The cumulative stock of EVs on China's roads surpassed 10 million by 2023, and Chinese manufacturers have achieved economies of scale that drive down costs, further propelling adoption. By contrast, **Europe as a whole saw a little over one in five new cars being electric in 2023** [21] – a very high figure by historical standards and up dramatically from just a few percent in mid-decade. Within Europe, certain countries lead: for instance, **Germany** reached a milestone in 2023 by joining China and the U.S. in the "million EVs sold per year" club [24], and countries like Norway (not an EU member but European) are near or above 80–90% EV for new sales, acting as harbingers of where others may head. On average, Europe's EV market share has stabilized around the 20% mark as some incentives were rolled back, but this still signals a market firmly in the steep uptake phase [25].

The **United States**, which trailed in EV adoption for many years, is now also climbing quickly. In 2023, **roughly one in ten new cars sold in the U.S. was electric** [21]. This is up from only about 2% in 2018 [26] – a five-fold increase in share in just five years – and the trajectory points sharply upward as more affordable models hit the market and charging infrastructure expands. The U.S. EV market crossed 1 million annual sales in 2023, making it the world's third-largest EV market by volume [22]. Still, the U.S. has considerable ground to cover to catch up to Europe and China in percentage terms, indicating that its EV S-curve started later and is a bit shallower so far. Reasons include cheaper gasoline, a love of large pickup trucks (only recently getting electric options), and initially weaker policy support – factors now changing with new federal incentives and dozens of new EV models available.

On the opposite end of the spectrum are advanced economies like **Japan**, and developing giants like **India**, where EV adoption remains in the very early stages – essentially the flat bottom of the S-curve. Japan, despite being a leading auto manufacturing nation and a pioneer in hybrids (like the Toyota Prius), has been surprisingly slow on pure EVs. As of **2022, only about 1–2% of new cars sold in Japan were fully electric** [27]. Japanese automakers and consumers have been cautious, preferring hybrid technology and constrained by limited charging infrastructure. In fact, one analysis noted that **Japan's BEV (battery electric vehicle) penetration in 2022 was just 2.2%, lower than even India's** [27]. Similarly, **India's EV market share was around 1.3% in 2022** [28] – meaning out of 3.8 million passenger vehicles sold in India that year, only about 50,000 were electric. While both Japan and India are poised for growth (India, for example, set ambitious targets of 30% of new car sales being electric by 2030 [29] [30]), for now they illustrate how the EV revolution's timing can vary widely. The **UK**, notably, is more aligned with the faster

European trend – the UK reached around **16% EV share by 2022** (and higher if plug-in hybrids are included), with plans to ban new gasoline car sales by 2030 which effectively forces the S-curve to completion by then.

When we consider **ramp-up data**, the curves are striking. It took until around **2015 for the world to sell its first million modern electric cars in a year** – a milestone built on early models like the Nissan Leaf, Tesla Model S, and Chevy Volt. But by 2021, annual EV sales exceeded 6 million [19]; by 2023, 14 million. The **cumulative stock of EVs on the road globally hit 40 million in 2023** [31], up from just 5 million or so in 2018. This five-year period was a turning point: **electric cars went from 2% of new sales in 2018 to 18% in 2023** [18], clearly entering the steep middle of the S-curve. Experts forecast that by mid-decade, the global share could surpass 25%, and by 2030 perhaps the majority of new cars sold worldwide will be electric – especially as major markets enforce stricter emissions rules. The ramp is further fueled by *economies of scale* in manufacturing (battery costs have plummeted more than 80% over the past decade) and *network effects* in infrastructure (more chargers beget more EV buyers, which beget more chargers, and so on).

However, even as EV sales charts shoot upward, one must remember that **fleet turnover in automotive is slow** – cars last 10-15 years – so the S-curve for the *installed base* of cars will lag behind sales. In the U.S., for example, EVs are ~10% of new sales, but under 1% of the 250 million cars on American roads. The true saturation point – when virtually all vehicles in use are electric – is likely a couple of decades away, not expected until the 2040s in leading regions (so long as new sales become nearly 100% electric by the 2030s). In any case, the past few years have decisively demonstrated an **inflection point for EV adoption**. The **traditional S-curve shape is evident**: after years of lingering in low single digits (despite plenty of hype and skepticism), EV sales hit a tipping point around 2020 in key markets, and the climb since has been extraordinary. Policy has played a role (for instance, the EU's CO_2 fleet targets, California's zero-emission vehicle mandate, China's subsidies and mandates), but so has consumer sentiment and word-of-mouth as EVs proved enjoyable and cheaper to run. Automakers themselves have pivoted from reluctance to a full embrace of electrification – many now announce end-dates for combustion engine development. As of 2023, we can say **the electric vehicle's transition from early adopter novelty to mainstream option is well underway**, even if the top of the curve (full saturation) is still on the horizon.

Industrial Robots: Automation's Global March

The presence of robots on factory floors is not new – industrial robots have been welding, painting, and assembling in automotive plants since at least the 1970s. Yet for a long time, their use was confined to specific sectors (notably car manufacturing) and a few leading countries (like Japan or Germany). The broader **adoption of industrial robots** has recently accelerated dramatically, revealing another S-curve as automation spreads across industries and regions. A key metric to track is **robot density** – the number of industrial robots per 10,000 manufacturing workers, which serves as an indicator of automation level. In 2015, the world average robot density was 69 per 10,000 workers [32]. By 2021, that global average more than doubled to **141 robots per 10,000 workers** [32]. This shows how rapidly automation has advanced in just six years. "Robot density is a key indicator of adoption," notes the International Federation of Robotics (IFR), and the latest data show a surging global average thanks to massive deployments in recent years [33].

Leading the charge has been **Asia, particularly China**. Asia's average robot density grew at an astonishing **18% compound annual growth rate from 2016 to 2021, reaching 156 units per 10,000 workers** [34]. For comparison, European industries saw an 8% annual growth in robot density in that period (to 129 units), and the Americas similarly 8% (to 117 units) [16]. Asia's rapid automation is largely the story of China's rise. **China has invested in industrial robotics on an unprecedented scale**, going from a relative laggard to the world's top adopter in a decade. In 2021, China overtook the United States in robot density for the first time [35]. China reached **322 robots per 10,000 manufacturing workers in 2021** [35], putting it fifth globally (behind only South Korea, Singapore, Japan, and Germany, which have long been highly automated) [36]. This marked an incredible milestone: **China, essentially a developing country in terms of automation a generation ago, is now among the world's most automated major industrial economies** [35]. The IFR attributed this to China's "massive investment" and noted that even at 322, China still has "much opportunity to automate," hinting the climb isn't over [33].

The raw installation figures tell an even starker tale of exponential growth. In 2021, Chinese factories installed a record **243,300 new industrial robots**, a **44% increase over the previous year** [37] [38]. To put that in perspective, China alone accounted for **about half of all industrial robot installations worldwide in 2021** [39]. This breakneck pace continued: in 2022, despite global economic headwinds, **China's robot installations grew another 5% to reach 290,258 units – again about 52% of the world's total for that year** [40]. The cumulative effect is staggering. The **operational stock of industrial robots in China surpassed 1 million units in 2021** and then hit **1.5 million units in 2022** [41]. China thus became the first country ever to have over a million robots in operation, and its stock was growing at **25% annually** in the five years up to 2022 [41]. For context, Europe's total stock of industrial robots in 2022 was about 728,000 units, and North America's was 452,000 [41]. China's robot population is not only the largest but expanding far faster, illustrating an S-curve on steroids as it catches up to and surpasses older industrial powers.

Of course, different countries are at different points on this curve. **Japan** was the original robotics powerhouse – in the 1980s and 90s, Japanese factories (in electronics and automotive) led the world in robot usage. Japan still ranks highly in robot density (it was third globally in 2021, after South Korea and Singapore [42]). However, Japan's adoption curve has matured; its density is high but growing more slowly now as it was an early adopter. **South Korea** stands out with an extraordinary 1,000 robots per 10,000 workers in 2021 (the highest in the world) [43], thanks to its dominant electronics and automotive sectors heavily automating – an example of a country near the top plateau of the S-curve. **Germany** likewise has steadily increased to over 300 per 10,000, reflecting a strong engineering sector [42]. The **United States** historically was slower to invest in robotics compared to some peers (relying more on cheaper labor or offshoring), but U.S. robot density has risen to around 274 per 10,000 by 2021 – now behind China's 322 [35]. The U.S. was actually surpassed by China in this metric, highlighting how the baton of rapid growth has passed to China. Yet American manufacturers are now rapidly catching up, with record orders for robots especially in sectors like logistics and fulfillment centers beyond the traditional automotive assembly use-case.

And then there's **India** – a country with a vast labor force and historically minimal robot usage. That too is beginning to change as wages rise and quality demands increase. In 2023, **India installed a record 8,510 industrial robots, up 59% from the previous year** [44]. While that number is tiny next to China's hundreds of thousands, it pushed India to rank 7th worldwide for new installations [44]. India's **operational stock of robots nearly doubled from 2018 to 2023, reaching about 45,000 units** [45]. Clearly, India is at the very bottom of its S-curve, just starting to accelerate. As Indian industry modernizes – especially in automotive, consumer goods, and electronics assembly – we can expect a steep climb in its automation adoption as well, albeit from a low base.

The implications of this global robotics ramp-up are significant. **Automation is no longer confined to a few wealthy countries; it's spreading everywhere**. The S-curve for industrial robots globally is in a high-growth phase: the world average density doubling in six years is proof of that [32]. We see a virtuous cycle where increased demand drives better technology and lower costs, which further spurs demand. For example, Chinese robot manufacturers have emerged to challenge the long-standing dominance of Japanese and European robot makers, bringing prices down. By 2022, **Chinese companies supplied 41% of the robots sold to China's electronics industry and 17% of those in its automotive industry** [46], undercutting foreign brands in cost. This competition and scale are driving a downward pressure on cost per robot, allowing more industries to justify automation. The IFR notes that the **global average robot density doubling since 2015** was fueled by such dynamics – as more units are produced, efficiencies improve and even mid-tier manufacturers adopt them [32].

It's also interesting to note that **the concentration of robot adoption is still high** – just five countries (China, Japan, the US, South Korea, and Germany) account for the bulk of global robot installations. But that concentration is slowly lessening as others join the fray. The IFR's 2022 report highlighted that China, Europe and the US together represent about two-thirds of total global robot stock (they also represent about two-thirds of manufacturing output, so it aligns) [47]. As these markets mature, the next wave might see countries like India, Brazil or Thailand picking up the pace. The S-curve of robotics might therefore have multiple waves as different economies automate on their own timeline.

In summary, **industrial robotics adoption has shifted into high gear globally in the last decade**. After years of incremental growth, the field hit an inflection point, especially visible in China's rise. The result has been year after year of record installation counts and swiftly climbing density rates. While some industries and regions are nearing a saturation point (e.g., South Korea's auto factories can hardly add many more robots without total human replacement, given their 1:10 human-to-robot ratio), others are just starting. We are, in effect, witnessing one of the largest transformations in manufacturing since the assembly line – a transformation measured in an S-curve of steel arms and automated guided vehicles populating factories from Bavaria to Bangalore.

Autonomous industrial robots assembling car bodies at a BMW factory. Advances in robotics have led to a surge in deployments on assembly lines worldwide, with robot density doubling globally between 2015 and 2021. [48] [33]

Personal Computers and the Internet: Laying the Digital Foundation

Before smartphones put a computer in every pocket, the **personal computer (PC) revolution** put one on every desk. The PC's adoption curve, though stretched over a longer period, exhibits the telltale S-shape as well. The first mass-market personal computers emerged in the late 1970s and early 1980s (the Apple II, Commodore PET, IBM PC, etc.), but for a while they were largely the domain of hobbyists, tech-savvy entrepreneurs, and certain business users. In 1984, only about 8% of U.S. households had a computer. Even by 1990, **about 15% of American households owned a personal computer** [49]. This was the slow initial phase – PCs were still relatively expensive and not yet seen as essential by most people. Then came the **1990s**, and with them a rapid escalation. As PCs became more affordable and useful (thanks to graphical interfaces like Windows 95 and applications like Microsoft Office and early internet connectivity), adoption **soared**. Between 1990 and 1997, the share of U.S. households with a computer jumped from 15% to **35%** [49] – more than doubling in seven years. Households across the socioeconomic spectrum started to see the value, especially for education and work. By 2000, roughly half of U.S. households had a PC, and by the early 2000s the majority did [50]. The U.S. Census Bureau noted that as of 2003, **62% of American households had one or more computers** [51], a number that only grew from there.

The trend continued into the 2000s until saturation. By **2016, nearly 89% of U.S. households had a computer at home** [52]. Essentially, the PC went from 0% to ~90% penetration in the U.S. over roughly 35 years, with the steepest climb in the 1990s. Today, the figure is in the mid-90s percentage (with the remaining households often being those of elderly individuals or in very low-income brackets). A similar S-curve played out in other advanced economies: for instance, the **United Kingdom's household computer ownership rose from negligible in the early 1980s to about 88% by mid-2010s** [53]. **Germany, France, Japan** – all saw PCs diffuse widely by the 2000s. Japan had a strong early computer market (NEC's PC-98, etc.), and by the 2000s Japanese households had high ownership rates, though some Japanese consumers eventually shifted more to mobile devices for personal use. One interesting variation is that **developing countries** in many cases leapfrogged directly to mobile computing without a phase of universal PC ownership. For example, in countries across Africa, Latin America, and parts of Asia, household PC penetration never reached the levels seen in the West, but mobile phone penetration (and later smartphone penetration) skyrocketed – essentially an alternate path up the digital adoption curve.

Nonetheless, the PC's importance cannot be overstated: it introduced the public to digital computing and paved the way for the internet. And speaking of the **internet**, its adoption might be one of the most consequential S-curves of all, transforming nearly every aspect of modern life. If we mark the internet's popular birth around the mid-1990s (after Tim Berners-Lee invented the World Wide Web protocol in 1989 [54], and the first browsers emerged in the early 90s), we see a very rapid ascent by historical standards. In **1995, only about 14% of U.S. adults had internet access** [55] (most via dial-up modems). Many Americans in the mid-90s had still never heard of the "information superhighway" – Pew surveys from 1994 found 42% hadn't heard of the internet at all, and another 21% were only vaguely aware of it [56]. But internet usage took off in the late 90s: by 2000 roughly half of U.S. adults were online. Indeed, the U.S. reached a majority of households online by 2001 [50]. After that, growth continued as broadband replaced dial-up and online services became more compelling. By 2010, about 70-75% of U.S. adults were internet users. And **as of 2019, around 90% of American adults used the internet at least occasionally** [57]. Current estimates put U.S. internet penetration (among the whole population) at about **91% in 2021** [58]. In other words, essentially everyone except some senior citizens or those in very remote areas is now connected. The U.S. internet adoption S-curve has largely plateaued in the high 80s to low 90s percent – the remaining gap may close slowly as holdouts diminish, but the era of double-digit annual growth in new users is long over.

Europe followed a similar pattern: early adoption in the Nordics and Western Europe in the 90s, mass uptake by the 2000s. Many European countries today also report around 90% of their population using the internet. **The UK** is over 90%, **Germany** around the same, **France** slightly lower but in the high 80s, etc. **Japan** likewise has very high internet usage rates (around 90%). In all these developed economies, the internet's S-curve was steep in the late 90s and 2000s, then leveled off as saturation neared. On the other hand, **China and India** illustrate the next wave of the internet adoption curve on a global scale. **China** had virtually no internet users in 1994; the few that existed were on academic networks. By the early 2000s, China had tens of millions of users as cyber cafés and home PCs with dial-up spread. The real boom was in the 2010s: cheap smartphones and China's homegrown internet giants (Tencent, Alibaba, Baidu, etc.) brought hundreds of millions online. By 2019, China had **over 850 million internet users** [9], making up about 60% of its population at that time. Today, China's internet penetration is around 70% and rising. It now boasts the largest online population of any country – over **1 billion users** – a remarkable climb from essentially zero just 25 years ago. **India**, too, saw a late 2010s explosion thanks to smartphones and ultra-cheap data. As of 2023, India likely has around 700 million internet users (roughly 50% penetration), up from maybe 100 million in 2010. India is on a steep part of its curve, with potentially hundreds of millions still to come online (many via mobile). This means the global internet adoption S-curve is still climbing, even though the Western world is near saturation.

To quantify the global picture: **in 2000, roughly 5-7% of the world's population was online** (mainly in rich countries). By 2010, about 30% of humanity had internet access. By 2019, it was 53%. And by 2024, the International Telecommunication Union (ITU) estimates **68% of the world's population – about 5.5 billion people – are using the internet** [59]. Put another way, in just the five years from 2019 to 2024, some **1.3 billion people came online for the first time** [59]. That's an extraordinary surge, driven heavily by mobile connectivity in developing regions. Still, **2.6 billion people remain offline as of 2024** [60], meaning the global S-curve has a ways to go before flattening out at universal access. Many of those offline are in rural parts of Africa and South Asia, or among marginalized groups, and will be the focus of connectivity efforts in the coming years. Nevertheless, the trend is clear: at the current pace, the internet user base might reach 90% of humanity by the 2030s, effectively completing one of the most rapid and consequential adoption curves in history.

It's important to highlight how intertwined the PC and internet curves were. The **1990s PC boom helped enable the internet boom**, since one needed a computer to go online in that era. Now, of course, smartphones have taken over as the primary on-ramp to the internet for billions. This means some late adopters skipped the PC stage entirely – a phenomenon known as "leapfrogging." For example, in sub-Saharan Africa, while only a minority of people own a personal computer, a majority have mobile phones and an increasing share have internet-enabled smartphones. So their **internet S-curve is being climbed via mobile** rather than the PC. This dynamic has made the internet adoption in developing regions often faster than it would have been if it depended on PCs and wired infrastructure.

In summarizing the digital foundational technologies: **the personal computer spread steadily from the 1980s, surged in the 1990s, and saturated in the 2010s in rich countries**, enabling the subsequent transformation of work and personal productivity. The **internet's adoption was even faster** – roughly 25-30 years from near-zero to two-thirds of the world – fundamentally altering commerce, communication, media and beyond. These S-curves laid the groundwork for everything from e-commerce to social media to cloud computing (which itself rode on the internet). They demonstrate how quickly a technology can become entrenched once value is evident and costs drop: what started as expensive tools for specialists (early PCs costing several thousand dollars, early internet requiring technical know-how to use) became utilities that nearly everyone on the planet expects to have access to.

Having examined these historical cases – smartphones, cloud, EVs, robots, PCs, and the internet – a clear pattern emerges of **accelerating adoption once a tipping point is reached**. In each case, the ramp-up involved interplay of innovation, falling costs, network effects, and often policy support. Now, we turn our attention to the **cutting-edge technologies of the present and near future**. Can we already see the first signs of new S-curves forming in the realms of **generative AI and autonomous robots**? How fast might these technologies scale relative to the benchmarks set by earlier innovations? And what factors will influence their trajectory? The second half of this report delves into these questions, exploring the early data and context for AI-based computer assistants and humanoid robots – two domains that are just beginning their journey from experimental to ubiquitous.

Generative AI Agents: A New Era of Rapid Software Adoption

Late one November night in 2022, a research lab called OpenAI quietly released a web interface for a conversational AI model named ChatGPT. Within days, social media was ablaze with examples of this chatbot's uncanny ability to answer questions, draft essays, and hold conversations. What followed was an **unprecedented surge in user adoption** for a new technology. In just two months, **ChatGPT reached 100 million monthly active users**, making it the fastest-growing consumer application in history at that time [61]. For comparison, it took TikTok about nine months to hit 100 million users, and Instagram about two and a half years [62]. ChatGPT's overnight success was a wake-up call: generative AI had arrived on the public stage, and there was a massive appetite for it. This explosive debut can be seen as the very beginning of an S-curve – one that might turn out to be **even steeper than past tech adoption waves, as some analysts predict** [63].

ChatGPT itself was just a precursor. The underlying technology – large language models and generative AI – is now being integrated into a wide array of products, heralding the era of **computer-using agents** and AI copilots. Unlike some past enterprise technologies that required building a user base from scratch, these AI agents are **piggybacking on existing platforms with huge user bases**, which could turbocharge their adoption. Take **Microsoft 365 Copilot**, for example. Announced in 2023, it is an AI assistant built atop OpenAI's GPT models, integrated across the Microsoft Office suite (Word, Excel, PowerPoint, Outlook, Teams, etc.). Microsoft 365 (formerly Office 365) is used by hundreds of millions of people worldwide – in fact, Microsoft reported about **382 million paid Office 365 seats as of early 2023** [64]. By embedding an AI agent into this suite, Microsoft essentially has a distribution channel to reach up to those hundreds of millions of users rapidly. Even if only a fraction initially enable or pay for the Copilot feature, that could still mean tens of millions of active users, essentially overnight once the feature is rolled out. Similarly, **Google is integrating generative AI ("Duet AI") into Google Workspace (Docs, Gmail, Sheets, etc.)**, which has its own enormous user base spanning enterprises, schools, and consumers using Gmail. In short, the moment these AI copilots become generally available, their potential reach is already global and massive.

This strategy of integration means **the adoption curve for AI agents could be extremely steep** – perhaps steeper than any previous enterprise software. Consider that previous productivity tools (like word processors or spreadsheets) had to spread gradually organization by organization, often through purchasing decisions and training. In contrast, an AI feature added to, say, Microsoft Word can appear for millions of users with a simple update. If Microsoft flips the switch to enable Copilot for all Microsoft 365 subscribers (or even a large subset), we could feasibly see a jump from near-zero to hundred-million-plus users in a matter of months. Even with a cautious rollout (perhaps starting with enterprise customers willing to pay a premium), the network effect and competitive pressure may drive a rapid expansion. Companies might feel **pressure to adopt AI copilots because their competitors are doing so**, creating a cascade. Early evidence of interest is strong: Microsoft's preview program for 365 Copilot reportedly had thousands of companies signing up to trial it. There's a recognition that these AI tools can boost productivity – from drafting emails to analyzing sales data – and no one wants to be left behind.

Moreover, the integration isn't limited to optional features. Microsoft is also building AI deeply into Windows itself. **Windows 11 introduced "Windows Copilot", effectively integrating a generative AI assistant at the operating system level**. This makes Windows 11 **the first PC platform to offer centralized AI assistance built-in** [65]. If AI becomes a default part of the OS that ships on new PCs, then every new computer could come with an AI agent ready to use. Suddenly, the adoption is not just rapid but near-automatic. Microsoft CEO Satya Nadella likened this moment to the introduction of the graphical user interface – a paradigm shift in how we interact with computers. If a billion Windows users all get an AI assistant in the next couple of years through OS updates, that dwarfs even ChatGPT's viral growth.

Likewise, **mobile ecosystems** are joining in. Both Apple and Google have been infusing AI into their mobile OS and services (think Siri, Google Assistant – though those are older and more limited forms of AI). We can expect more advanced generative AI in those assistants soon. In effect, the adoption of **AI-based agents might ride on the back of smartphone adoption**. And since there are over 5.5 billion smartphone users globally [66] , integration at the OS or app level on mobile could bring hundreds of millions more into contact with AI agents without them explicitly seeking it out.

Another dimension is **enterprise integration beyond Office apps**. Many companies are embedding generative AI into their workflows: customer service chatbots powered by GPT, software development copilots (e.g., GitHub Copilot, which uses AI to assist programmers), decision-support tools in ERPs and CRMs (Salesforce has rolled out Einstein GPT for CRM). These domain-specific agents might not boast the headline user numbers of Office or Windows, but they signify a broad penetration of AI into daily business operations. For instance, GitHub Copilot was launched in late 2021; by mid-2022 it already had over 1 million users and was writing significant portions of code for those developers. Microsoft even reported that in some programming tasks, **Copilot could generate 40% or more of the code** – an astonishing uptake in a short time (anecdotally, many developers now consider AI assistance a standard part of their coding toolkit). The point is, **the product-market fit for generative AI in many knowledge tasks is being found very quickly**, and once found, scaling is largely a matter of cloud deployment rather than physical distribution.

That said, there are gating factors to consider. Unlike physical products, **scaling software AI has minimal manufacturing constraints**, but it does face **computational and cost constraints**. Running large AI models for hundreds of millions of users requires vast computing power (data centers, GPUs, etc.). Companies like Microsoft and Google are investing heavily in AI supercomputers to meet this demand, but the rollout might be throttled by infrastructure in the short term (for example, initially limiting how many queries a user can run, or which users get access). Privacy and security are another factor – enterprises will test these AI agents in limited settings before fully deploying, to ensure they don't leak data or make egregious errors. Despite these, the overall trajectory seems inevitable: **AI copilots are on track to become as commonplace as the applications they reside in**.

Adoption speed will also depend on **user trust and behavior change**. Some employees may resist or underutilize AI features initially, out of habit or concern. But if history is any guide (think of how quickly email or search engines became indispensable), once people see colleagues using AI to work faster or smarter, they will adopt it too. In a sense, we might soon consider it unthinkable to compose a complex document or analyze data without an AI helper – much as it became unthinkable to do business without email or to do research without Google.

We are still in the **early days of the generative AI S-curve**, perhaps analogous to where the internet was in, say, the mid-1990s or where smartphones were around 2008. There's excitement, experimentation, and also plenty of hype. But the empirical signs of a true inflection are visible. Already, **by mid-2023 there were estimates that well over 100 million people worldwide had tried some form of generative AI tool** (ChatGPT or equivalents). Some surveys suggest even higher numbers if we include those who used AI-generated content indirectly. **UBS analysts wrote in early 2023 that they "cannot recall a faster ramp in**

a consumer internet app" than ChatGPT [67], and many commentators noted how generative AI seemed to achieve mainstream awareness faster than even social media did back in its day.

One reason to believe **this adoption could outpace prior technologies** is the **relative ease of deployment**. To get the benefits of electricity, for example, homes had to be wired and appliances bought – a slow process. To get the benefit of AI, often one just needs to click a feature in software one already has, or log into a web service. The barrier to entry is low; often the service might even be free or included in existing plans. Another reason is **AI's generality** – it's not a single-purpose tool but a multi-purpose assistant. This broad applicability can drive adoption across many domains at once (writing, coding, marketing, data analysis, customer support, etc.), creating a multiplier effect on user growth.

There is also **a flipside risk**: if generative AI fails to meet expectations (for instance, if early corporate deployments produce a lot of incorrect outputs or even scandals like leaked sensitive info), there could be a temporary pullback – a "trough of disillusionment" – before the technology improves. This would resemble other hype cycles where initial excitement is followed by a reality check, then a slower rise to true productivity. Indeed, some executives are cautious: concerns about factual accuracy of AI ("hallucinations") and security mean some companies are going slow. But these challenges are actively being addressed with model improvements and guardrails, and the overall competitive drive to use AI for efficiency likely outweighs the hesitations in the long run.

Perhaps the most illustrative comparison is that generative AI might follow a path **similar to the internet itself**: a burst of hype, some bubbles (e.g., countless AI startups now vying for attention), but underneath that, a real transformation that steadily expands and eventually becomes ubiquitous. As one venture capitalist observed, generative AI is broad and abstract, a platform like the internet was, and **while it might take a decade to figure out all the best use cases, the eventual adoption could be vast** [68]. The difference is that the **time from invention to mass deployment** could be much shorter for AI. The internet took 30+ years from its early academic days to reach 3 billion users; smartphones took about 15 years from the first modern smartphone to reach 3+ billion users. It's not inconceivable that **AI assistants could reach billions of users in well under a decade**, given they can propagate through existing devices and apps.

In summary, the stage is set for generative AI agents to sprint up the adoption curve. The early spike with ChatGPT shows raw demand; the impending integration into software and systems we already use will supply the means. If earlier technological diffusion seemed like a marathon, this might shape up to be a 100-meter dash. And while it's always wise to temper predictions (no technology is adopted overnight in every corner of the world), many experts indeed argue that the **adoption of generative AI may outpace any prior innovation in scope and speed** [63]. We are, to use a Silicon Valley cliché, likely "still day one" for this technology – but day two may come very fast.

Humanoid Robots: Lab Experiments to Factory Helpers

Amid the boom in virtual AI agents, another frontier of technology is materializing – quite literally in physical form: **humanoid robots**. Long a staple of science fiction and aspirational R&D projects, human-like robots that can walk, manipulate objects, and perform general tasks have never been produced at scale. That may be about to change. In recent years, there's been a flurry of activity in this space, suggesting we may be at the dawn of an S-curve for **humanoid robot adoption**. The trajectory here will be very different from software: it involves hardware manufacturing, safety certification, and physical jobs. But the **first**

signs of real-world deployment are now emerging, and industry leaders are openly talking in terms of millions – even billions – of units in the future [69].

A landmark moment was in 2021 when **Tesla announced its "Optimus" humanoid robot project**, leveraging its expertise in batteries, motors, and AI. Skepticism was high initially – many saw it as a moonshot diversion. But by 2022 and 2023, Tesla had developed prototype units of Optimus and was showcasing them performing basic tasks (albeit often via tele-operation or scripted motions). Elon Musk has repeatedly emphasized that he sees Optimus as a product with potentially enormous scale, even more significant than Tesla's car business in the long run. At Tesla's 2024 investor event "We, Robot," Musk gave concrete (if ambitious) numbers: he predicted that **Optimus could cost under $20,000 and be produced in quantities eventually reaching millions per year** [70]. In early 2025, Musk went further on a conference call: after noting Tesla's plan to build about 10,000 robots in 2025 as an initial target, he stated **"it won't be long before Tesla is making 100 million of these things a year"** [71]. He even projected that by the 2040s there could be **"a billion humanoid robots on Earth"** [69]. These jaw-dropping figures, if realized, would dwarf the automotive industry (which produces around 85 million vehicles a year globally). Of course, such claims are coming from an avowed optimist (and salesman) for the technology, so they should be taken with caution. But they underscore a vision of a future where **humanoid robots might become as commonplace as cars or smartphones are today**.

Translating that vision into reality, however, depends on near-term progress and proving the use cases. The initial strategy for deploying humanoid robots is to focus on controlled environments like factories and warehouses – places that are structured, yet still designed for humans to move around and perform tasks. A major development in this vein occurred in January 2024: **Robotics startup Figure announced a partnership with BMW Manufacturing to test its humanoid robots in a BMW car factory in South Carolina** [72]. This was the **first known commercial deal to put general-purpose humanoids into a real production environment**. The plan was to start with a small number of units, evaluate their performance on tasks like moving materials or working in the body shop, and then potentially expand if things go well [73]. Figure's CEO noted this as huge validation – after all, BMW's factories are highly optimized, and any automation they introduce must justify itself in efficiency. The Figure robots are slated to be integrated over 12–24 months, implying that by 2025 or 2026 we could see humanoids side by side with human workers on some assembly lines [74].

Other carmakers have also dabbled or signaled interest. **Honda** has a long history with humanoid robots (the ASIMO robot developed in the 2000s), though it was more a demonstration than a factory tool. **Hyundai**, after acquiring Boston Dynamics in 2020, has been exploring humanoid and legged robots for industrial use [69]. In fact, Boston Dynamics' bipedal robots (like Atlas) have shown remarkable agility in demos, but they remain expensive prototypes. **Tesla**, for its part, plans to first use Optimus robots within its own factories – doing simple repetitive jobs like carrying parts or perhaps assembly tasks that haven't been worth building a fixed robot for. In late 2023, Tesla reported it had a few Optimus prototypes "working" in a factory (though evidence was scant) and stated a goal to have several thousand doing useful work by end of 2025 [75]. If they achieve that, it would mark the beginning of real productivity from humanoids, and likely embolden wider adoption.

The journey to **product-market fit for humanoid robots** is just starting. The key is finding tasks that these robots can do reliably and cost-effectively that either couldn't be automated before or where replacing a human has clear benefits (due to labor shortage, cost, or safety). Early tasks might include things like: intralogistics (moving parts or carts around a warehouse/factory), machine tending (loading/unloading machines), simple assembly or inspection tasks, or working in hazardous environments. The advantage of a humanoid form factor is that it can, in theory, operate in spaces designed for humans and use human tools and interfaces without those environments being rebuilt. If Optimus or Figure can demonstrate, say, that one robot can save the labor of one human on a loading dock and operate for less cost than a salary, that would be a strong economic case. Given the high cost of labor in some markets and difficulties in hiring for certain repetitive jobs, many industries are watching closely.

Once a compelling use case is demonstrated, adoption could ramp up fast – within the limits of manufacturing capacity. Factories, warehouses, retail chains, hospitality (imagine hotel cleaning or room service robots), eldercare – the potential markets are vast, but the requirements differ. It's likely we'll see a phase where supply is actually the bottleneck: demand outstrips how many robots can be produced. Indeed, Musk's commentary about ramping to 100k per month production lines [76] suggests they anticipate massive demand if they can get the robots to work properly. Tesla revealed it is designing dedicated assembly lines for robots: first 1,000 units/month, then scaling to 10,000/month, and then a Version 2 design at 100,000/month capacity [76]. These numbers, while hypothetical for now, give a sense of how scaling might occur in stages: a pilot line, a medium-scale line, then a giga-factory for robots. **If Tesla or others achieve a production rate of 1 million robots per year, it would be unprecedented in robotics history** – no other general-purpose robot has been made at that scale (industrial robot arms total global production is a few hundred thousand per year, for example).

Manufacturing ramp-up brings in the question of **supply chain and cost reduction**. Building humanoid robots at scale will heavily rely on sourcing key components (motors, actuators, sensors) at low cost. This is where global supply chains and competition play a pivotal role. Analysts have pointed out that **Tesla's goal of a <$20,000 robot is heavily dependent on leveraging China's manufacturing ecosystem** [77]. Many core components like precision actuators, gearboxes, and bearings are currently made most cost-effectively in China [78]. Chinese suppliers have years of experience and intense competition that have driven down prices for things like robot servos and ball screws [79]. If trade tensions or export controls cut off access to these, it could impede cost reduction. For now, it seems likely that any mass-produced robot will incorporate a lot of components from China or similarly efficient manufacturing hubs. As volumes increase, we'd expect costs to follow a **learning curve** (Wright's Law) where each doubling of cumulative production might reduce unit cost by, say, a certain percentage. In electric cars, for instance, Tesla managed significant cost declines through scaling; they'd aim to do similar for robots. Early units might cost tens of thousands to build, but by the time you've made 100k of them, maybe cost halves, and by a million, halves again, etc. **Scale and competition will put strong downward pressure on prices** – as more companies (and there are several humanoid startups now: Figure, Tesla, Sanctuary, Apptronik, Agility Robotics, etc.) enter, each will look for ways to cut costs to gain market share. The SCMP report underlined that **without access to cost-efficient suppliers, mass production at target prices might be "put on hold"** [77], emphasizing how critical the global supply chain is to this S-curve. But assuming no catastrophic trade disruptions, the likely scenario is similar to other electronics: design in the West, build at scale where it's cheapest, drive costs down, thereby unlocking more demand – a positive feedback loop.

So what might the adoption curve for humanoid robots look like? Initially, very slow – we are basically at year zero. Perhaps a few dozen in 2023 (mostly prototypes), maybe hundreds in 2024-25 in pilot programs, then thousands by 2026-27 if pilots succeed. After that, if the technology has proven value, it could take off in the 2030s. It's plausible that by the early 2030s, annual production could be in the hundreds of thousands globally (Tesla alone is aiming for that within a decade). Reaching **100 million units in operation** might happen by late 2030s or 2040s if things really accelerate – that would require production scaling into millions per year and sustained for a couple of decades. Musk's vision of **a billion robots by 2040s** [69] implies essentially an S-curve that hits the steep part in the 2030s and then saturates at a level where perhaps there is roughly one humanoid robot per 10 humans on Earth. That saturation level is speculative – it assumes they become widely useful not just in factories but in **commercial settings, public service, and even private homes**. Indeed, for a billion robots, we're talking about something that would likely serve individual households (as helpers or caregivers) or be pervasive in society (like robotic assistants in every store, hospital, and transit station).

There are historical analogies: the automobile took about 50-60 years to go from invention to one car per several persons in wealthy countries, and now about 1.4 billion vehicles are in use worldwide (roughly one per six people on Earth). If humanoid robots address enough needs, perhaps they could follow a somewhat similar trajectory but maybe faster given a more tech-accelerated age. Still, hardware tends to have a slower adoption than pure software due to production limits. Cars needed roads, fuel, repair networks; robots will need maintenance, software updates, etc. There will also be societal and regulatory questions: safety standards for robots working around people, public acceptance (people might be uncomfortable initially with humanoids in everyday life, requiring a cultural acclimation), and labor impact (could spark pushback or regulation if they start displacing many jobs). Those factors can either slow or shape the path of adoption.

In factories and businesses, where ROI dictates decisions, if robots save money, they'll be adopted quickly. BMW's experiment with Figure robots is telling – if BMW finds even moderate success, other manufacturers will race not to fall behind. Supply chain issues aside, demand could become frenzied in specific sectors facing labor shortages (e.g., countries like Japan with aging populations might eagerly adopt caregiving robots for elder care if they work well; Amazon might deploy legions of robots in warehouses if they can handle the work). One could imagine **network effects** of a sort: as more companies use robots, best practices emerge, the robots get better (their AI learns from more data), and costs drop, which encourages even more use.

By the late 2020s, we should have a clearer sense of the humanoid robot S-curve shape. It might still be modest – say tens of thousands in use – or it might be clearly trending exponentially. **Investor interest is already strong**, with significant capital flowing into startups aiming to build these robots, indicating a belief that a large market will materialize. The renewed interest noted by Reuters, where investors see humanoids as having potential because they can "learn to perform new tasks like humans do" [80], reflects an expectation that these robots won't be static machines but continually improving agents, much like software.

In conclusion, **the emergence of humanoid robots stands at a similar juncture as personal computers did in the late 1970s or EVs in the early 2010s** – a technology with long development behind it, now nearing viability for real use, but not yet widely deployed. If the prototypes can transition to practical workers, the adoption curve could follow the classic S-shape: slow initial adoption (few early users in 2020s), then accelerating growth once cost and performance hit the right threshold, and eventually a plateau when most organizations that need a robot have one (and perhaps households too, if they become consumer devices). The timeframe could be compressed relative to older technologies thanks to better manufacturing techniques and global connectivity of knowledge. But it will still involve physical scaling, which tends to run in years, not weeks.

As of now, we can foresee the **first significant influx of humanoids in factories by around 2025-2026**, broader industrial adoption through the late 2020s, and possible expansion to public and home environments in the 2030s if all goes well. The milestone targets mentioned – 100 million units, 1 billion units – are ambitious and perhaps speculative, but not impossible if we stretch the horizon to mid-century. Musk's prediction of a billion by 2040s [69] might be optimistic, but it's a useful provocateur: it forces us to imagine a world where **humans and human-like machines coexist at scale**, analogous to how we coexist with billions of computers and smartphones today. And much like those earlier tech adoptions, cost declines from scale and competition will be key. It is telling that experts already emphasize that the $20k price target for Optimus hinges on tapping into China's mature, cost-effective manufacturing base [77] [78] – a hint that the classical forces of globalization and economies of scale will drive this S-curve as surely as they did for smartphones or solar panels (another technology that saw a steep adoption as costs fell).

Conclusion: Toward the Next Technological Plateaus

Looking back over the last few decades, the deployment of smartphones, cloud services, electric vehicles, industrial robots, personal computers, and the internet each illustrates the power of the S-curve as a narrative of technological change. In every case, early growth was deceptively slow – sometimes leading pundits to declare the technology had "failed to take off" – but then one day the conditions were right and adoption exploded. **Empirical data from multiple regions confirm this pattern**: the U.S., EU, UK, China, Japan, and India all experienced these S-curves, though with different timing and slopes. The U.S. saw rapid PC and internet uptake in the 90s, China caught up in mobile phones and now leads in EVs and robots, India is late but accelerating in smartphones and internet, and so on. These regional contrasts show that technology often spreads unevenly at first (often concentrated in wealthier or more ready markets), but over time tends to globalize – the S-curve eventually encompasses the world, not just its starting point.

Now, as we stand in the mid-2020s, **we are witnessing the beginning of new S-curves in generative AI and humanoid robotics**. The narrative nonfiction of technological change is entering a new chapter. The second half of this decade will likely see whether AI agents become as common as web browsers and whether humanoid robots become as normal in factories as industrial arms are today. Early signs are exceedingly promising for AI agents – integration into ubiquitous software could propel adoption faster than any prior tool. As one venture capital article noted, **once a technology hits a tipping point, things "develop very quickly"** [3], and generative AI may be reaching that point now, with broader adoption potentially "accelerating even more quickly than past waves" [63]. For humanoid robots, the tipping point might still be a few years off until the product is refined, but given the momentum (major companies testing them, significant capital investment, and clear labor needs), their ramp-up could surprise us in speed once it truly begins.

It's also crucial to recognize the **synergy between these trends**. Advanced AI will likely make robots more capable (AI brains for mechanical bodies), and widespread robotics will create demand for even better AI to control them. Enterprises might adopt AI software and robotic automation together as part of a general digital transformation wave. Consider a future factory where AI copilots help design products and manage logistics, while humanoid robots handle physical assembly – all coordinated via cloud systems. The S-curves of different technologies can reinforce one another. Cloud computing's growth, for example, has enabled the AI revolution by providing the needed computing power on demand; smartphone adoption has driven internet usage higher, and so on.

One common thread in all these adoption curves is **cost and accessibility**. Technologies hit mass adoption when they become affordable and easy to use. PCs became mainstream when costs fell and interfaces improved. The internet needed cheap access and user-friendly web browsers. Smartphones took off once competitive manufacturing (largely in Asia) made them inexpensive enough for the masses. EVs are now taking off as battery costs drop and range anxiety eases with more chargers. For AI agents, cost is less about purchase price (often they come as a service) but more about **trust and utility** – once they demonstrably save time or money, they'll be indispensable. For humanoid robots, cost is currently a huge barrier (a prototype might cost hundreds of thousands of dollars to develop), but scale manufacturing aims to bring it to the price of a car or lower, at which point businesses can justify the investment widely.

Competition plays a key role in accelerating adoption too. When multiple firms vie in the same space, they tend to innovate and reduce prices faster. We saw that with smartphone makers, with cloud providers, with EV manufacturers – and we are seeing it now with AI startups and soon with robot makers. This competition not only drives down costs but also pushes the technology to be better and more tailored to users, thus attracting more adopters and reinforcing the S-curve's steep ascent.

Finally, history teaches us that as technologies saturate, they often become **invisible infrastructure** – taken for granted. Few today marvel at a smartphone's existence; it's assumed. The internet is like electricity, just there. We may be headed that way with AI agents – they might soon be as mundane as Clippy once was (hopefully more useful!), just an ever-present assistant we don't even think about. Humanoid robots, being physical and evocative, might not become "invisible" so much as normalized – future generations may find it unremarkable to see a robot cleaning the streets or stocking shelves at the store.

The pace of technological adoption is unquestionably accelerating in aggregate, but each technology still has to navigate its own path with human, economic, and regulatory factors. The case studies of smartphones, cloud, EVs, robots, PCs, and the internet show that when the time is ripe – a confluence of innovation maturity, cost reduction, and demand – the scaling can be breathtakingly swift. Generative AI has already demonstrated a flash of that potential, and humanoid robots might follow suit in the coming decade. The exact timelines are hard to predict, but if there is one lesson from the past, it is to not underestimate the speed of adoption once an S-curve hits its stride.

Today's **advanced AI copilots and nascent humanoids may seem novel, even experimental**, much as the early web or early mobile phones once did. But if we project current trends forward, backed by evidence and the economic logic of scale, it's plausible that by the 2030s we'll talk about them in the same matter-of-fact way we discuss cloud storage or smartphone apps. The diffusion of innovation marches on, and we find ourselves at the exciting inflection points of new curves – living through the very narrative of technological transformation that future analysts will study. As always, the transition may feel slow until, suddenly, it is astonishingly fast. And by the time these technologies reach their plateaus, our world – work, society, daily life – will likely have been reshaped in ways we can barely begin to imagine, just as a pre-internet or pre-smartphone world is hard to imagine today.

Sources:

- B2Venture – *"When Hype Cycles Meet S-Curves"*, on generational tech adoption patterns [1] [12] [63].
- Pew Research – U.S. smartphone ownership rise from 35% in 2011 to 91% by 2023 [6].
- Wikipedia (2021 data) – Smartphone penetration ~72% US, ~60% China, ~43% India [9].

- Datacenter Post/Gartner – Virtualization >50% of server workloads by 2014, projected 86% by 2016 [13].
- B2Venture – Cloud adoption 5% of workloads (2015) to ~60% today [12].
- IEA Global EV Outlook 2024 – EVs were 2% of global new cars in 2018, ~18% in 2023 [18]; one-third of new cars in China, one-fifth in Europe, one-tenth in U.S. in 2023 [21]; limited sales in Japan/India [21].
- India Briefing – India EVs just 1.3% of 2022 car sales [28].
- IFR Press (2022) – China robot density 322 per 10k workers, surpassing U.S. [35]; global average density 141 in 2021 vs 69 in 2015 [32].
- BusinessWire/IFR (2023) – China installed 290k robots in 2022 (52% of world) [40]; China's stock 1.5M vs Europe 728k, NA 452k [41].
- Reuters – BMW partnering to deploy Figure humanoid robots in a U.S. plant [72]; Elon Musk's prediction of a billion humanoid robots in 2040s [69].
- Electrek (Jan 31, 2025) – Musk aiming for 10k Optimus in 2025 and "100 million per year" eventually [71]; < $20k price once at 1M/year production [81].
- SCMP (Apr 2025) – Experts say $20k cost depends on China's supply chain; Chinese competition drives down component prices [77] [78] [79].
- Reuters – ChatGPT reached 100M users in 2 months, fastest-ever adoption [61].
- Wikipedia – As of 2019, 90% of U.S. adults use the internet [57]; Statista – U.S. internet usage ~91% in 2021 [58]; ITU – 68% of world online in 2024 (5.5B users) vs 53% in 2019 [59].

---

1 2 3 4 12 17 63 66 68 b2venture
https://www.b2venture.vc/stories/when-hype-cycles-meet-s-curves-the-roll-out-conundrum-of-generative-ai

5 10 Smartphone Usage Statistics 2025 (Worldwide Data) - Demand Sage
https://www.demandsage.com/smartphone-usage-statistics/

6 Demographics of Mobile Device Ownership and Adoption in the United States
https://www.pewresearch.org/internet/fact-sheet/mobile/

7 8 9 List of countries by smartphone penetration - Wikipedia
https://en.wikipedia.org/wiki/List_of_countries_by_smartphone_penetration

11 Global smartphone penetration 2016-2024 - Statista
https://www.statista.com/statistics/203734/global-smartphone-penetration-per-capita-since-2005/

13 14 15 Virtualization has Surpassed 50 percent of all Server Workloads | Data Center POST
https://datacenterpost.com/virtualization-surpassed-50-percent-server-workloads/

16 32 33 34 35 36 42 43 China overtakes USA in robot density - International Federation of Robotics
https://ifr.org/ifr-press-releases/news/china-overtakes-usa-in-robot-density

18 19 20 21 22 23 24 26 31 47 Trends in electric cars – Global EV Outlook 2024 – Analysis - IEA
https://www.iea.org/reports/global-ev-outlook-2024/trends-in-electric-cars

25 [PDF] China, Europe, United States, and India, 2023 H1
https://theicct.org/wp-content/uploads/2024/01/ID-87-%E2%80%93-Maj-Mkts-H1_final2.pdf

27 Japan: battery electric vehicle market share 2024 - Statista
https://www.statista.com/statistics/711994/japan-electric-car-market-share/

28 29 30 India’s EV Production Capacity and Domestic Auto Market Trends
https://www.india-briefing.com/news/indias-prospects-as-an-ev-hub-consumer-market-and-production-capacity-30157.html/

37 China: Robot installations grew by 44 percent
https://ifr.org/ifr-press-releases/news/china-robot-installations-grew-by-44-percent

38 IFR: China's robot installations rose by 44% in 2021
https://www.therobotreport.com/ifr-chinas-robot-installations-rose-by-44-in-2021/

39 China Bought 44% More Robots in 2021 Than It Did in 2020, Finds ...
https://www.robotics247.com/article/china_bought_44_more_robots_in_2021_than_it_did_in_2020_finds_ifr_report

40 41 46 Record 1.5 million Robots Work in China's Factories – IFR reports
https://www.businesswire.com/news/home/20230925370684/en/Record-1.5-million-Robots-Work-in-Chinas-Factories-IFR-reports

44 India Ranks 7th in Annual Robot Installations Worldwide – IFR reports
https://www.roboticstomorrow.com/news/2024/09/24/india-ranks-7th-in-annual-robot-installations-worldwide-%E2%80%93-ifr-reports/23208

45 International Federation of Robotics 's Post - LinkedIn
https://www.linkedin.com/posts/international-federation-of-robotics_2024-sep-24ifrpressreleaseworldrobotics-activity-7247086647392010240--Ns4

48 69 72 73 74 80 BMW taps humanoid startup Figure to take on Tesla's robot | Reuters
https://www.reuters.com/business/autos-transportation/bmw-taps-humanoid-startup-figure-take-teslas-robot-2024-01-18/

49 Computer ownership up sharply in the 1990s : The Economics Daily : U.S. Bureau of Labor Statistics
https://www.bls.gov/opub/ted/1999/apr/wk1/art01.htm

50 57 Internet in the United States - Wikipedia
https://en.wikipedia.org/wiki/Internet_in_the_United_States

51 [PDF] Computer and Internet Use in the United States: 2003
https://www.census.gov/content/dam/Census/library/publications/2005/demo/p23-208.pdf

52 U.S. households with PC/computer at home 2016 - Statista
https://www.statista.com/statistics/214641/household-adoption-rate-of-computer-in-the-us-since-1997/

53 UK households: ownership of home computers 1985-2018 - Statista
https://www.statista.com/statistics/289191/household-penetration-of-home-computers-in-the-uk/

54 55 56 Part 1: How the internet has woven itself into American life | Pew Research Center
https://www.pewresearch.org/internet/2014/02/27/part-1-how-the-internet-has-woven-itself-into-american-life/

58 United States internet penetration 2000-2025 - Statista
https://www.statista.com/statistics/209117/us-internet-penetration/

59 60 Statistics
https://www.itu.int/en/itu-d/statistics/pages/stat/default.aspx

61 62 67 ChatGPT sets record for fastest-growing user base - analyst note | Reuters
https://www.reuters.com/technology/chatgpt-sets-record-fastest-growing-user-base-analyst-note-2023-02-01/

64 Office 365 Reaches 400 Million Users
https://office365itpros.com/2024/01/31/office-365-reaches-400-million/

65 Bringing the power of AI to Windows 11 - unlocking a new era of ...
https://blogs.windows.com/windowsdeveloper/2023/05/23/bringing-the-power-of-ai-to-windows-11-unlocking-a-new-era-of-productivity-for-customers-and-developers-with-windows-copilot-and-dev-home/

70 71 75 76 81 Elon Musk says Tesla aims to build 10000 Optimus robots this year
https://electrek.co/2025/01/31/elon-musk-says-tesla-aims-to-build-10000-optimus-robots-this-year/

77 78 79 Tesla can’t make Optimus robot for US$20,000 without China, humanoid experts say | South China Morning Post
https://www.scmp.com/economy/global-economy/article/3307085/tesla-cant-make-optimus-robot-us20000-without-china-humanoid-experts-say